



**MSc Dissertation: Predicting Employee Attrition and Deriving
Strategic Retention Insights: A Business Analytics Approach
Using Machine Learning**

Abstract

Employee attrition represents a persistent challenge for organisations, costing high rates of recruitment, lost productivity, and institutional knowledge loss. This dissertation examines how **predictive analytics** can model voluntary turnover while aligning with organisational governance, managerial credibility, and financial prudence.

Six machine learning models **Logistic Regression, Support Vector Machine, Decision Tree, Random Forest, XGBoost**, and **Naive Bayes** were benchmarked on a HR dataset (around 16% attrition prevalence) using a leakage-safe pipeline. Pre-processing consisted of median/mode imputation, one-hot encoding with unknown handling, selective scaling, and n-fold SMOTE for class imbalance handling. Models were evaluated by cross-validation and a held-out test set using ROC-AUC, PR-AUC, calibration checks, and threshold sensitivity analysis, recognising that the default 0.50 cut-off may not be suitable for organisational priorities. A cost-sensitive framework translated predictive performance into approximate financial outcomes.

Results show that linear models prefer recall at the cost of reduced precision, while ensemble methods (Random Forest, XGBoost) achieving stronger precision–recall trade-offs. Threshold optimisation had a stronger effect of false positives and false negatives with a direct effect on HR workload and financial cost. At a governance-aligned operating point (precision floor ≥ 0.60), the Random Forest achieved PR-AUC at 0.525 and ROC-AUC at 0.87 flagging 7.8% of employees to generate majority-positive, capacity-aligned intervention lists. SHAP analyses revealed the three principal drivers like Workload strain (**OverTime**), Organisational embeddedness (**YearsAtCompany**; **YearsWithCurrManager**), and perceived alternatives (**MonthlyIncome**, **JobLevel**, **StockOptionLevel**) moderated by **JobRole** and **BusinessTravel**.

This dissertation contributes to both scholarship and practice by incorporating cost analysis, threshold governance, and fairness audits into attrition prediction, going beyond accuracy metrics to decision relevance. Recommendations do include incorporating threshold-aware policies, supporting support-first interventions, and embedding quarterly monitoring to ensure calibration, fairness, and model credibility. The study demonstrates how predictive models can be built to be technically robust, ethically deployable, and relevant with managerial decision-making.

Acknowledgements

I would like to express my deepest gratitude to my supervisor, **Professor Konstantinos Katsikopoulos** for their expert guidance, constructive feedback, and continuous encouragement throughout the development of my dissertation. Their insightful comments and patient support have been invaluable in improving my ideas and ensuring the academic rigour of this work.

I would like to thank the faculty and staff at the University of Southampton Business School for providing us the resources, training material and academic environment which enabled me to undertake this dissertation.

This study has benefited greatly from access to the **IBM HR Analytics Employee Attrition** dataset, which has provided myself with the foundation for creating and assessing the predictive models. I acknowledge the dataset providers for making this resource publicly available for academic research.

I would like to express my sincere gratitude to my family and friends for their constant encouragement, understanding, and moral support during this journey. Their patience and belief in me have been a constant source of motivation for this dissertation.

Finally, I wish to extend my appreciation to my peers and colleagues for their helpful discussions, shared resources, and constructive feedback, all of which enriched my research experience.

Contents

Chapter 1: Introduction	8
1.1 Background and Context.....	8
1.2 Problem Statement	8
1.3 Motivation.....	9
1.4 Research Aim and Objectives	9
1.5 Research Questions and Hypotheses.....	10
1.6 Theoretical Focus	11
1.7 Scope and Boundaries.....	11
1.8 Methodological Overview.....	11
1.9 Significance of Research.....	12
1.10 Structure of Dissertation	12
Chapter 2: Literature review	14
2.1 Chapter Aim.....	14
2.2 Review Approach.....	15
2.3 Theoretical Perspectives on Turnover.....	15
2.3.1 Attitudinal and Behavioural Models	16
2.3.2 Unfolding and Shock Models	16
2.3.3 Resource-Based Perspectives.....	16
2.3.4 Embeddedness Theory	16
2.3.5 Economic Perspectives	17
2.3.6 Critical Synthesis	17
2.4 Predictive HR Analytics.....	17
2.5 Evaluation and Calibration under Class Imbalance	20
2.6 Analytical Method Description	23
2.6.1 Model Benchmarking.....	23
2.6.2 Class Imbalance Handling	23
2.7 Summary of Literature	24
2.8 Research Gaps.....	25
2.9 Positioning for This Dissertation	25
Chapter 3 – Methodology	26
3.1 Introduction.....	26
3.2 Research Philosophy and Approach.....	26
3.3 Research Design.....	27
3.4 Data Source and Features.....	28

3.5 Preprocessing Pipeline	30
3.6 Modelling Strategy.....	31
3.7 Evaluation Framework.....	33
3.8 Threshold Governance	34
3.9 Interpretability.....	35
3.10 Fairness Audits.....	36
3.11 Governance Artefacts	37
3.12 Ethical Considerations	38
3.13 Limitations of Methodology	38
3.14 Chapter Summary	39
Chapter 4 - Results and Analysis	40
4.1 Introduction.....	40
4.2 Test-Set Evaluation	41
4.3 Confusion Matrices at Threshold = 0.50 (Baseline)	46
4.4 Threshold = 0.40 (Recall-Oriented Sensitivity Sweep)	51
4.5 Governance Threshold = 0.60 (Finalist RF)	53
4.6 Feature Importance and Interpretability.....	55
4.7 Cross-Validation as Robustness	56
4.8 Reliability, Limitations, and Robustness Checks.....	58
4.9 Synthesis	60
Chapter 5 – Discussion and Findings.....	62
5.1 Strengths and Weaknesses of the Research.....	64
5.2 Personal Learning Journey.....	65
5.3 Fresh Insights for Research and Practice	65
5.4 Future Research Directions	66
5.5 Conclusion to the Discussion.....	66
Chapter 6 – Conclusion and Recommendations	68
6.1 Introduction.....	68
6.2 Summary of Key Findings	69
6.3 Theoretical Contributions	70
6.4 Practical Recommendations.....	70
6.5 Limitations of the Study.....	71
6.6 Directions for Future Research	72
6.7 Final Reflections	73
7. References.....	74

List of Tables

1. **Table 1.** Comparative strengths of model families.
2. **Table 2.** Evaluation framework for model assessment, distinguishing threshold-free metrics, thresholded performance, sensitivity analyses, and governance-level decision add-ons.
3. **Table 3.** Predictor categories and theoretical links.
4. **Table 4.** Model families and rationale.
5. **Table 5.** Evaluation metrics and decision relevance.
6. **Table 6.** Threshold governance scenarios.
7. **Table 7.** Fairness audit structure.
8. **Table 8.** Test-set evaluation metrics (Accuracy, Precision, Recall, F1, ROC-AUC) across six machine learning models.
9. **Table 9.** Confusion matrix results at baseline threshold ($T = 0.50$) across six machine learning models.
10. **Table 10.** Governance threshold outcomes ($T = 0.60$) across six machine learning models, summarising ROC-AUC, precision, recall, and governance decision criteria.
11. **Table 11.** Cross-validation results (mean $\pm$ standard deviation) across six machine learning models, comparing ROC-AUC, F1, precision, and recall stability.

List of Figures

1. **Figure 1.** Methodology flowchart – End-to-end pipeline from data preprocessing to governance.
2. **Figure 2.** ROC curve for all six models.
3. **Figure 3.** Comparative bar chart of test-set model performance (Accuracy, Precision, Recall, F1, ROC-AUC).

4. **Figure 4.** Confusion matrix at baseline threshold ($T = 0.50$) across six machine learning models.
5. **Figure 5.** Precision–recall trade-off curves at thresholds 0.50 and 0.40 across six models, illustrating the recall gains and precision losses from lowering the decision boundary.
6. **Figure 6.** Random Forest feature importance ranking (Top 10 predictors of employee attrition).

Chapter 1: Introduction

1.1 Background and Context

In modern-day organisations, voluntary employee attrition has emerged as one of the most critical strategic challenges (Wright & McMahan, 1992). As compared to involuntary turnover, which can often be controlled through managerial decisions, voluntary exits represent an employee's conscious choice to leave. Such departures create a ripple effect on costs — the visible cost of recruitment and training, the loss of tacit knowledge, the weakening of customer and team relationships, and the morale shocks that often trigger further resignations (Hom et al., 2017; Hausknecht & Trevor, 2011). Replacement of an individual employee is estimated to cost between 150% and 200% of the yearly wage (Portfolio, 2025; Deloitte, 2017) and in knowledge-based industries, the loss of a few important staff can destabilise the entire business units. Current studies further show retention rate is increasingly associated along with strategic HRM practices, while firms that integrate HR strategy with business goals achieve stronger outcomes in repression of turnover (Impact of Strategic Human Resource Management Practices on Employee Retention, 2024).

Despite all these high-level stakes, many organisations remain reliant on reactive retention solutions such as exit interviews, trailed employee surveys, or managerial intuition. These approaches surface signals after employees have already disengaged, leaving the managers with little or no room to intervene. Predictive analytics therefore offers a much proactive alternative with the potential to identify at-risk employees earlier and to align interventions with organisational capacity. Yet the promise of prediction alone is insufficient - unless models are accurate, interpretable, calibrated, fair, and governable, they risk being disregarded, distrusted, or even harmful in use (Bassi, 2011).

1.2 Problem Statement

Existing approaches to attrition prediction frequently fall short of bridging the gap between technical modelling and organisational decision-making. Many studies report headline accuracy or AUC scores but neglect issues of calibration, class imbalance or threshold sensitivity. In practice, a model that identifies all stayers will achieve deceptively high accuracy

in datasets where attrition prevalence is low, but it delivers no value to managers seeking to retain employees.

Further, attrition models carry reputational and ethical risks. Predictors such as commute distance or overtime may inadvertently encode demographic biases, exposing organisations to fairness concerns. Without robust governance artefacts such as model cards, monitoring reports, and fairness audits, predictive HR tools risk being dismissed as “black boxes” with little legitimacy.

The central research problem is therefore not whether attrition can be predicted, it can but how predictive models can be designed, validated, and governed to be technically robust, ethically deployable, and practically usable for strategic HRM.

1.3 Motivation

This research is motivated by both academic and managerial imperatives.

From an academic perspective, voluntary attrition represents a fertile case through which to test how modern machine learning (ML) can extend and be disciplined by turnover theory. Classical frameworks such as Human Capital Theory, Equity Theory, Job Embeddedness Theory, and the Job Demands–Resources model identify clear antecedents of attrition. Yet these constructs are rarely operationalised through HRIS data and predictive pipelines that managers can use in practice. This project provides an opportunity to close that gap by translating theory-rich insights into calibrated, explainable, and governance-ready models.

From a managerial perspective, HR functions face a mounting pressure to adopt evidence-based retention strategies. Employee churn remains a top concern in annual surveys of executives, yet most firms lack proactive analytical tools (Bassi, 2011; Deloitte, 2017). This project demonstrates how HR data that organisations already collect (tenure, workload, compensation) can be harnessed responsibly to reduce attrition risks.

1.4 Research Aim and Objectives

The overarching aim of this dissertation is to design, validate, and evaluate a predictive modelling pipeline for employee attrition that balances accuracy, interpretability, fairness, and

governance.

This aim is operationalised through four objectives:

1. **Benchmarking models:** Compare six modelling families - Logistic Regression, Linear SVM, Decision Tree, Random Forest, XGBoost, and Gaussian Naïve Bayes under class imbalance, selecting a finalist based on PR-AUC.
2. **Threshold governance:** Operationalise a precision-floor policy (≥ 0.60) by selecting an operating threshold T from out-of-fold predictions, maximising recall within that constraint.
3. **Interpretability:** Provide global and local SHAP explanations that link predictions directly to managerial levers such as rota smoothing, progression planning, and mentoring.
4. **Fairness and governance:** Audit precision and recall across Gender and Age groups and formalise monitoring artefacts such as model cards, threshold memos, and challenger models.

1.5 Research Questions and Hypotheses

To achieve these objectives, the dissertation addresses the following questions:

- **RQ1.** Which modelling family performs best under class imbalance for tabular HR attrition?
- **RQ2.** What operating threshold maximises recall subject to a precision floor?
- **RQ3.** Which features drive predicted risk, and how can explanations guide managerial interventions?
- **RQ4.** How consistent is the performance across subgroups and what mitigations are appropriate?

Hypothesis: Ensemble methods (Random Forest, XGBoost) will outperform linear models on PR-AUC, but Random Forest will always provide the best balance of precision, calibration, and interpretability for practical HR adoption.

1.6 Theoretical Focus

The research is based on established turnover theories, which together provide the conceptual foundation for selecting predictors, designing models, and interpreting results:

- **Workload strain:** Longitudinal demands such as overtime and travel reflect resource depletion, consistent with burnout and resource-depletion perspectives.
- **Organisational embeddedness:** Tenure and Manager–employee ties reflect relational anchors, essential to embeddedness perspectives on turnover.
- **Perceived rewards and alternatives :** Income, job level and stock options mapped onto opportunity structures, aligning well with equity and human capital perceptions.
- **Contextual moderators:** JobRole, Department, and Work–Life Balance embody how strain, embeddedness, and alternatives interact, as posited with contingency perspectives.

This theoretical grounding helps in shaping of methodological design: linear models provides understandable baselines, ensemble models capture non-linear interactions, and SHAP explanations relocates outputs to theory, ensuring that predictions remain actionable for managers (Wright & McMahan, 1992; Hom et al., 2017).

1.7 Scope and Boundaries

The research adopts a bounded case study using secondary, de-identified HR data from a single firm with an attrition prevalence of around 16%. The dataset includes demographic, workload, compensation, and tenure-related variables but excludes qualitative constructs such as leadership style or organisational culture. Analysis is cross-sectional rather than longitudinal, limiting causal inference. Fairness audits were restricted to Gender and Age because of data availability. These boundaries provide reproducibility and transparency and also acknowledges limitations (Hausknecht & Trevor, 2011; Deloitte, 2017).

1.8 Methodological Overview

The study adopts a quantitative, positivist stance and a deductive design. A leak-safe pipeline is employed end-to-end: median/mode imputation → one-hot encoding → scaling for linear

models → SMOTE confined to cross-validation folds → estimator training. All models are trained under identical conditions using established Python libraries.

Evaluation prioritises PR-AUC to reflect class imbalance, supported by ROC-AUC, calibration reliability curves, and threshold sensitivity analysis. Interpretability is achieved through SHAP explanations, while fairness is assessed through subgroup audits and governance artefacts (Bassi, 2011; Deloitte, 2017).

1.9 Significance of Research

This study is significant on three levels:

- **Theoretical contribution:** It extends extant turnover theories through empirical validation of their constructs in HRIS data and reprioritising underexplored drivers such as workload intensity.
- **Methodological contribution:** It demonstrates that modern ML techniques like calibration, PR-AUC, threshold policies, and SHAP explanations can be embedded in HR analytics to bridge the gap between accuracy reporting and decision aid-gap.
- **Practical contribution:** It provides HR managers with a governance-ready early-warning system that is calibrated, fair, and transparent, making it possible for interventions to be proactive without overburdening resources or eroding trust.

1.10 Structure of Dissertation

The dissertation follows a problem-to-practice outline:

- **Chapter 2** critically surveys literature on turnover determinants, predictive HR analytics, evaluation under imbalance, and fairness/governance, inspiring the methodological design.
- **Chapter 3** delineates the dataset, preprocessing pipeline, six-model benchmark, validation protocol, threshold selection, interpretability techniques and subgroup auditing.
- **Chapter 4** reports on model performance, threshold sensitivity, calibration, SHAP explanations, fairness audit and robustness checks.
- **Chapter 5** describes the findings, proposing a retention and governance strategies along reflection on limitations and personal learning.

- **Chapter 6** concludes by addressing the research questions, outlining contributions to theory and practice, and outlining opportunities for future research

Chapter 2: Literature review

2.1 Chapter Aim

Voluntary employee attrition has long been recognised as one of the most consequential issues in organisational management and research. Turnover drains human and social capital, disrupts customer and team relationships, and increases operational costs through recruitment, training, and onboarding (Price, 1977; Price, 2001; Griffeth, Hom & Gaertner, 2000; Allen, Bryant & Vardaman, 2010). Industry studies estimate that replacing an employee typically costs between 150 and 200 per cent of their annual salary, without accounting for intangible losses such as morale contagion and weakened culture (Allen, Bryant & Vardaman, 2010; Batt & Colvin, 2011). In knowledge-intensive industries, the departure of a few highly skilled employees can destabilise entire functions, making attrition not only an HR concern but a strategic risk multiplier (March & Simon, 1958; Price, 2001).

The academic study of attrition has evolved significantly over the past six decades. Early work in organisational behaviour sought to explain why people leave, emphasising job satisfaction, organisational commitment, and attitudinal predictors (Porter & Steers, 1973; Mobley, 1977; Herzberg, Mausner & Snyderman, 1959; Vroom, 1964). These explanatory models provided rich theoretical insights but often lacked predictive precision and timeliness. With the digitisation of human resource information systems (HRIS), attention has shifted toward predicting who is most likely to leave, enabling proactive interventions rather than reactive responses (Hausknecht & Trevor, 2011; Davenport, Harris & Shapiro, 2010; Angrave et al., 2016). This evolution parallels broader movements in evidence-based management and the rise of predictive analytics in other domains such as finance, healthcare, and marketing (Davenport, Harris & Shapiro, 2010; Angrave et al., 2016).

The purpose of this chapter is to critically assess the foundation of attrition studies and predictive HR analytics, providing a theoretical and methodological basis for the modelling pipeline used in this dissertation. The review is divided into two broad sections for this first half: (i) theory perspectives on turnover, which lays out the central constructs and debates that underpin voluntary attrition, and (ii) predictive HR analytics, which considers the shift away from explanatory models to machine learning and the methodological concerns that remain.

2.2 Review Approach

This review adopts a systematised narrative approach, combining rigour in search and inclusion with flexibility for critical synthesis (Angrave et al., 2016). A structured database search was performed across Scopus, Web of Science, and Google Scholar, using keyword sets that combined employee turnover, voluntary attrition, job satisfaction, job embeddedness, JD-R, predictive HR analytics, class imbalance, calibration, and algorithmic fairness.

Inclusion criteria:

- Empirical peer-reviewed and methodological studies published between 2010 and 2024.
- Seminal texts as the formation for theoretical underpinning of turnover research, including March & Simon's Organizations, Equity Theory, and Job Embeddedness (March & Simon, 1958; Adams, 1963; Mitchell et al., 2001).
- Tabular HRIS transference-worthy research.

Exclusion criteria:

- Purely perceptual studies reliant solely on self-reported surveys, unless they provided transferable constructs (e.g., fairness, embeddedness).
- Grey literature and practitioner blogs lacking academic rigour.

In addition, reference snowballing was used to identify influential contributions in fairness, governance, and predictive modelling, such as Model Cards and internal algorithmic auditing (Mitchell et al., 2019; Raji et al., 2020). The aim was not to achieve exhaustive coverage but rather a discerning synthesis. This means interrogating what research converges on and diverges from and critically, where they fall short in providing actionable and governance-relevant predictive tools.

2.3 Theoretical Perspectives on Turnover

Turnover has been examined through several complementary yet sometimes contradictory traditions. Collectively, these traditions identify three broad mechanisms: strain and resource depletion, organisational embeddedness, and perceived alternatives and fairness.

2.3.1 Attitudinal and Behavioural Models

The earliest explanations emphasised attitudes. Job dissatisfaction and weakened organisational commitment were seen as precursors to withdrawal cognitions and exits (Porter & Steers, 1973; Mobley, 1977; Herzberg, Mausner & Snyderman, 1959; Vroom, 1964). Griffeth, Hom & Gaertner's (2000) meta-analysis considered that low satisfaction predicts turnover across industries. These models are helpful in highlighting the psychological origins of attrition but have the limitation of being survey-based and incapable of accommodating time sequencing. They tell us why people may consider leaving but not when or who will leave (Deci & Ryan, 1985; Kahn, 1990).

2.3.2 Unfolding and Shock Models

Recognising that not all exits follow a gradual dissatisfaction path, Lee & Mitchell (1994) proposed the unfolding model. They recognized the role of shocks - unsolicited offers, organisational conflicts, or personal life events that trigger re-evaluation and sudden exits. Maertz, Stevens & Campion (2007) showed shocks that interact with embeddedness and perceived alternatives, explaining sudden departures. Shocks are captured in HRIS thus weakening their predictive power. For prescriptive analytics, chronic signals (e.g., workload, tenure, pay trajectories) are more actionable (Rubenstein et al., 2018).

2.3.3 Resource-Based Perspectives

Resource theories define attrition as a consequence of resource depletion. Excessive demands (OverTime, BusinessTravel) drain individual capacity, while few resources accelerate burnout (Maslach & Leiter, 1997; Kahn, 1990). Rubenstein et al. (2018) proved workload strain consistently predicts attrition across sectors. Yet effects vary by career stage and context, necessitating interaction-aware modelling (e.g., OverTime $\times$ Tenure).

2.3.4 Embeddedness Theory

Mitchell et al. (2001) shifted explanations to why people stay, introducing job embeddedness: the links, fit, and sacrifices that bind employees. Tenure, manager retention, and department stability are HRIS proxies. Embeddedness protects against shocks (Halbesleben & Wheeler,

2008), yet tenure effects may be non-linear, with hazard spikes during managerial transitions (Rubenstein et al., 2018). Over-embedding can even keep disgruntled employees trapped, reducing but not preventing turnover. This is why there is a need of non-linear, flexible models such as Random Forests (Huang & Gartner, 2020; Sharma & Mishra, 2021).

2.3.5 Economic Perspectives

Economic theories do bring in rational and comparative elements. Equity Theory emphasises equity, arguing employees evaluate their pay relative to others (Adams, 1963). Plateauing growth is a forecaster of attrition than fixed remuneration, and fairness tends to be a more effective motivator than mere salary levels (Griffeth, Hom & Gaertner, 2000; Rubenstein et al., 2018). Expectancy Theory includes motivation by linking effort and performance with reward congruence (Vroom, 1964; Deci & Ryan, 1985). Classic organisational perspectives also frame mobility decisions as part of broader intra-organisational equilibria (March & Simon, 1958; Price, 1977; Price, 2001).

2.3.6 Critical Synthesis

Despite differences, turnover theories converge on three HRIS-measurable mechanisms: strain, embeddedness, and alternatives. However, there are still arguments on the relative significance of absolute versus relative pay, the protective vs hazardous standing of tenure, and the predictability of shocks. These contradictions justify predictive methods that can capture non-linearities and interactions, reconciling competing effects (Mitchell et al., 2001; Rubenstein et al., 2018).

Implication (RQ3): HRIS features - overtime, tenure, pay progression, manager continuity provides operational proxies for the turnover theories. Yet actionable insights require the advanced models capable of capturing complexity and translating patterns into managerial levers (Allen, Bryant & Vardaman, 2010; Batt & Colvin, 2011).

2.4 Predictive HR Analytics

The digitisation of HR data has enabled a methodological shift from explanatory regressions to predictive machine learning. Traditional logistic regression remains widely used for its

interpretability and well-calibrated probabilities, while attrition data often exhibiting multicollinearity, non-linearity, interaction effects that linear models can't capture effectively (Davenport, Harris & Shapiro, 2010; Angrave et al., 2016).

Machine learning expands the modelling toolkit:

Naïve Bayes is computationally efficient but relies on unrealistic independence assumptions (Huang & Gartner, 2020).

Support Vector Machines (SVMs) provide robust classification boundaries but require calibration for probability estimation (Sharma & Mishra, 2021).

Decision Trees are interpretable and simple to understand but get unstable and subject to overfitting (Sharma & Mishra, 2021).

Random Forests (RF) consists of multiple trees, offering accuracy, robustness and strong practical interpretability through feature importance (Huang & Gartner, 2020; Sharma & Mishra, 2021).

XGBoost (Extreme Gradient Boosting) delivers accuracy through boosting but is sensitive to tuning and are less transparent (Huang & Gartner, 2020).

Model family	Handles interactions	Calibration (native)	Data-prep sensitivity	Deployment relevance
Logistic Regression	Limited	Often strong	Scaling + one-hot encoding	Transparent baseline; good governance challenger (Angrave et al., 2016; Davenport, Harris & Shapiro, 2010)
Linear SVM	Linear only	Requires post-hoc scaling	Scaling + one-hot encoding	Robust margins; less interpretable (Sharma & Mishra, 2021)
Gaussian Naïve Bayes	Very limited	Often mis specified	Sensitive to correlation	Lightweight baseline; sanity check (Huang & Gartner, 2020)

Decision Tree	Yes	Variable; prone to overfit	Low	Illustrative but unstable (Sharma & Mishra, 2021)
Random Forest	Yes, robust	Good post-hoc	Low	Stable; SHAP-compatible; reproducible in practice (Huang & Gartner, 2020; Sharma & Mishra, 2021)
XGBoost / GBM	Yes, rich	Often needs post-hoc	Moderate	High accuracy; tuning-sensitive (Huang & Gartner, 2020)

Table 1: Comparative strengths of model families

Empirical case studies support these dynamics. Huang & Gartner (2020) predicted gradient boosting to HR attrition, achieving high accuracy but also raising interpretability concerns. Sharma & Mishra (2021) tested neural networks but concluded that Random Forests outperformed on tabular HR data. Such evidence motivates the construction of ensemble methods that weigh interpretability against accuracy.

However, benchmarking practices in HR analytics remain inconsistent. Many studies compare only two families (e.g., LR vs DT) or fail to standardise preprocessing, calibration, and imbalance handling; such inconsistencies undermine comparability and external validity (Angrave et al., 2016; Davenport, Harris & Shapiro, 2010). Leak-prone validation also inflates reported performance, misleading practitioners (Angrave et al., 2016).

Why Random Forest Emerges as Finalist: Among ensemble methods, Random Forest balances predictive power, stability, and interpretability in ensemble methods. It is more reproducible in governance use than gradient boosting because it is less sensitive to hyperparameters. Its simplicity of use with post-hoc explanation allows for predictions to be translated into managerial action (Huang & Gartner, 2020; Sharma & Mishra, 2021; Mitchell et al., 2019).

Implication (RQ1): This dissertation benchmarks six families - LR, SVM, NB, DT, RF, and XGB under a uniform, leak-safe pipeline, anticipating Random Forest as the most credible operational finalist while retaining Logistic Regression as a calibrated challenger (Angrave et al., 2016; Davenport, Harris & Shapiro, 2010).

2.5 Evaluation and Calibration under Class Imbalance

A central challenge in attrition modelling is class imbalance. In most HRIS datasets, only 10–20 per cent of employees exit voluntarily in a given year, meaning the positive (attrition) class is a minority. This creates distortions: a naïve model predicting “all stayers” can achieve 80–90 per cent accuracy, while missing all leavers. Such models may appear statistically strong but provide no managerial value (Allen, Bryant & Vardaman, 2010; Batt & Colvin, 2011).

To overcome this, researchers emphasise performance metrics sensitive to imbalance. While ROC-AUC is widely used, Precision–Recall AUC (PR-AUC) often better reflects organisational decision value, as it balances precision (the credibility of flagged cases) with recall (the coverage of actual leavers) (Huang & Gartner, 2020; Sharma & Mishra, 2021). For HR, recall is especially critical: the cost of a false negative (missing a leaver) is typically far higher than a false positive (a precautionary conversation with a stayer) (Allen, Bryant & Vardaman, 2010).

Calibration is another neglected but essential dimension. Managers act on probabilities, not ranks. If a model assigns an employee a 0.70 attrition risk, roughly 70 per cent of similar employees should indeed leave. Most ML models generate mis-calibrated probabilities undermining trust even when there is robust AUC; reliability curves and Brier scores provide calibration diagnostics (Huang & Gartner, 2020; Sharma & Mishra, 2021). Miscalibration has the potential to cause HR decision-makers to lose faith in model output, lowering adoption (Angrave et al., 2016).

Threshold choice adds another layer of complexity. Most studies default to 0.50 as a cut-off, but this is arbitrary. In HR contexts, thresholding must be treated as an explicit policy decision reflecting cost asymmetry. Missing a preventable leaver is often far more damaging than engaging unnecessarily with a stayer. Therefore, this dissertation enforces a precision-floor rule (≥ 0.60): ensuring that at least 60 per cent of flagged employees are genuine leavers, while maximising recall subject to that floor. The optimal threshold is frozen pre-test to avoid post

hoc bias, and sensitivity envelopes at 0.60/0.65/0.70 are reported to aid governance discussions (Allen, Bryant & Vardaman, 2010; Mitchell et al., 2019).

Alternative frameworks thus provide further depth to the toolkit. Cost-sensitive classifiers explicitly weight false negatives more heavily than false positives (Sharma & Ghosh, 2022). Decision curve-style analyses translate predictive outputs into expected utility, providing managers with clear trade-off visualisations (Allen, Bryant & Vardaman, 2010). Additional performance metrics such as F1, G-mean, and Matthews Correlation Coefficient (MCC) provide complementary checks. Yet in HR practice, simplicity matters: PR-AUC, calibration checks, and precision-floor thresholding balance rigour with interpretability (Angrave et al., 2016; Davenport, Harris & Shapiro, 2010).

Implication (RQ2): Adding PR-AUC as the primary metric, probability calibration, and threshold operation places assessment in organisational context, making it reproducible and decision-relevant (Mitchell et al., 2019; Raji et al., 2020).

Evaluation type	Metrics	Interpretation
Threshold-free	PR-AUC, ROC-AUC, calibration	Compare models fairly; assess probability quality (Huang & Gartner, 2020; Sharma & Mishra, 2021)
Thresholded (at T*)	Precision, Recall, F1, Flag-rate	Trade-off between workload and missed leavers (Allen, Bryant & Vardaman, 2010)
Sensitivity analysis	Floors (0.60 / 0.65 / 0.70)	Audit stricter policies; workload vs trust trade-offs (Mitchell et al., 2019)
Governance add-ons	Decision curves; cost-based analysis	Link operating point directly to business value (Raji et al., 2020)

Table 2: Evaluation framework for model assessment, distinguishing threshold-free metrics, thresholded performance, sensitivity analyses, and governance-level decision add-ons.

While predictive HR analytics is of transformative potential, it may also reproduce or even reinforce discrimination. HRIS features often act as proxies for protected characteristics: commute distance may correlate with age, income growth with gender, and overtime with caregiving responsibilities. Even if models don't contain protected attributes, indirect discrimination can still emerge (Barocas & Selbst, 2016).

The fairness literature is complex. Kleinberg, Mullainathan & Raghavan (2017) proved the impossibility theorem: under varying subgroup base rates, no model can satisfy all fairness criteria (e.g., equal precision, recall, and calibration for groups) at once. This means fairness is not an endpoint but a trade-off requiring governance. Mehrabi et al. (2021) advocate parity envelopes acceptable ranges of disparity as more pragmatic than strict guarantees.

Governance artefacts operationalise these principles. Model Cards document the purpose, dataset, features, metrics, and limitations of a model (Mitchell et al., 2019). Threshold Memos formalises policy choices, such as the 0.60 precision floor used in this dissertation. Monitoring Packs track PR-AUC, calibration, and subgroup parity quarterly, providing evidence of accountability. Challenger Models (e.g., Logistic Regression) run in parallel to find drift and provide rollback safeguards (Raji et al., 2020).

Ethical and regulatory frameworks reinforce the need for governance. Broader AI ethics guidance add to that AI must be fair, transparent, and accountable (Floridi et al., 2018; Lepri et al., 2018; Leicht-Deobald et al., 2019). Industry failures illustrate the risks of neglecting governance: widely reported hiring algorithms were abandoned after perpetuating bias, damaging trust and reputation (Holstein et al., 2019). Practitioners are more likely to adopt ML when transparency artefacts are in place (Mitchell et al., 2019; Raji et al., 2020).

Implication (RQ4): Predictive attrition models must embed subgroup audits (e.g., Gender, Age) and governance artefacts to ensure fairness, accountability, and organisational legitimacy. Without these, adoption will falter, regardless of technical accuracy (Barocas & Selbst, 2016; Kleinberg, Mullainathan & Raghavan, 2017; Mehrabi et al., 2021).

2.6 Analytical Method Description

The analytical design of this dissertation synthesises theoretical insights and methodological best practices into a governance-ready pipeline. Each choice is informed by the gaps seen in the literature.

2.6.1 Model Benchmarking

Six families - Logistic Regression, Linear SVM, Gaussian Naïve Bayes, Decision Tree, Random Forest, and XGBoost are benchmarked under uniform preprocessing. Logistic Regression provide an open baseline while SVM and Naïve Bayes are light comparators; Decision Trees exhibit instability; Random Forest offers stability and interpretability and XGBoost tests state-of-the-art boosting (Huang & Gartner, 2020; Sharma & Mishra, 2021). Benchmarking across these families offsets fragmented comparisons in previous HR research and practice (Angrave et al., 2016; Davenport, Harris & Shapiro, 2010).

2.6.2 Class Imbalance Handling

Resampling is applied within cross-validation folds so as to avoid leakage into the test set and ensure minority (leaver) cases are adequately represented (Huang & Gartner, 2020; Sharma & Mishra, 2021). Prior HR analytics often ignored imbalance, leading to inflated accuracy (Angrave et al., 2016).

Evaluation Metrics:

PR-AUC has precedence, with ROC-AUC and calibration reliability curves as secondary checks (Huang & Gartner, 2020; Sharma & Mishra, 2021). This aligns with organisational needs: balancing credibility (precision) with coverage (recall) (Allen, Bryant & Vardaman, 2010).

Threshold Governance:

Rather than using arbitrary defaults, a precision-floor rule (≥ 0.60) is enforced. The threshold T is frozen pre-test to ensure reproducibility, and sensitivity envelopes are reported. This operationalises thresholding as a policy choice (Mitchell et al., 2019; Raji et al., 2020).

Interpretability:

Post-hoc explanations are used to connect predictions to turnover theory. Global patterns pick

up systematic drivers (e.g., overtime as a sign of strain), local set of explanations accounts for individual cases (Mitchell et al., 2019; Holstein et al., 2019).

Fairness Audits:

Precision and recall are checked across Gender and Age subgroups, reported as parity envelopes. Fairness issues raised through predictive HR analytics are solved by it (Barocas & Selbst, 2016; Mehrabi et al., 2021).

Governance Artefacts:

Model Cards, Threshold Memos, Monitoring Packs, and Challenger Models create accountability and transparency (Mitchell et al., 2019; Raji et al., 2020). This renders the pipeline deployable in real-world HR environments.

Governance-ready attrition pipeline

Turnover theory → HRIS features → uniform benchmarking → PR-AUC + calibration → precision-floor threshold → explanations → fairness audits → governance artefacts.

Analytical Contribution:

This design portrays the four major gaps identified: fragmented benchmarking, superficial evaluation, theory-to-action translation, and weak governance. It produces a pipeline that is not only well accurate but also equally calibrated, interpretable, fair, and governable (Angrave et al., 2016; Davenport, Harris & Shapiro, 2010; Mitchell et al., 2019; Raji et al., 2020).

2.7 Summary of Literature

- Turnover theories converge on strain, embeddedness, and alternatives but disagree on relative weightings and boundary conditions (Porter & Steers, 1973; Mitchell et al., 2001; Rubenstein et al., 2018).
- Predictive HR analytics expands the methodological toolkit but suffers from inconsistent benchmarking and calibration neglect (Angrave et al., 2016; Davenport, Harris & Shapiro, 2010).
- Evaluation practices over-rely on accuracy and ROC, failing to prioritise PR-AUC, calibration, and threshold governance (Huang & Gartner, 2020; Sharma & Mishra, 2021;

Mitchell et al., 2019).

- Governance and fairness are fairly acknowledged but not always deployed in HR analytics, leaving the models vulnerable to bias and mistrust (Barocas & Selbst, 2016; Kleinberg, Mullainathan & Raghavan, 2017; Mehrabi et al., 2021).

2.8 Research Gaps

- 1. Theory-to-practice translation:** Robust constructs like strain and embeddedness rarely have actionable levers. Post-hoc explanations and governance artefacts provide this missing link (Mitchell et al., 2019; Holstein et al., 2019).
- 2. Benchmarking limitations:** HR analytics often compares narrow subsets under inconsistent pipelines. A leak-safe, multi-family benchmark addresses this (Angrave et al., 2016).
- 3. Evaluation blind spots:** All studies report accuracy and ROC while ignoring PR-AUC, calibration, and threshold governance. This study embeds all three (Huang & Gartner, 2020; Sharma & Mishra, 2021; Mitchell et al., 2019).
- 4. Fairness and governance practice:** Despite consensus on threats, few papers did implement artefacts such as model cards or challenger model (Mitchell et al., 2019; Raji et al., 2020; Floridi et al., 2018).

2.9 Positioning for This Dissertation

This dissertation directly addresses these gaps by:

- **RQ1:** Benchmarking six model families under a uniform, leak-safe pipeline (Angrave et al., 2016; Davenport, Harris & Shapiro, 2010).
- **RQ2:** Embedding a frozen precision-floor threshold policy (≥ 0.60) (Mitchell et al., 2019).
- **RQ3:** Translate turnover theory into self-explanatory action mapping (Holstein et al., 2019; Allen, Bryant & Vardaman, 2010).
- **RQ4:** Embedding fairness audits and governance artefacts (Barocas & Selbst, 2016; Raji et al., 2020; Floridi et al., 2018).

By integrating theory, predictive analytics, evaluation, and governance, this dissertation reframes attrition prediction as a multidimensional governance challenge. The goal is not simply to build an accurate model but to create a calibrated, fair, and actionable decision-

support system capable of efficiently catching more leavers while maintaining organisational legitimacy and trust (Mitchell et al., 2019; Raji et al., 2020).

Chapter 3 – Methodology

3.1 Introduction

Every dissertation must make explicit not only what was done but why it was done in that way. This chapter therefore sets out the methodological framework through which the research questions introduced in Chapter 1 are addressed and the theoretical debates reviewed in Chapter 2 are operationalised. The aim is not to provide a purely technical manual of procedures, but to offer a coherent narrative that links philosophical stance, research design, dataset selection, analytical pipeline, and governance considerations into one rigorous methodological architecture.

The chapter unfolds in stages. It starts off by explaining the research's philosophical stance, situating the project within a positivist, deductive tradition that prefers quantitative analysis. It then discusses the choice of research design: a bounded case study of secondary HRIS data, not framed as a convenience but as a theoretically informed decision in keeping with the study's aims. Following this, the chapter describes the dataset, its scope, features, and preparation, explicitly linking variables to turnover theories such as strain, embeddedness, and alternatives. The later sections (3.5–3.13) develop the analytical pipeline, model benchmarking, evaluation metrics, threshold governance, interpretability, fairness audits, governance artefacts, ethics, and limitations.

3.2 Research Philosophy and Approach

Research methodology is underpinned by assumptions about the nature of knowledge (ontology) and how it can be acquired (epistemology). This study assumes a positivist ontology: it assumes that employee attrition patterns exist independently of the researcher and can be identified through systematic analysis of data. It also embraces a deductive epistemology, beginning from established theories of turnover (e.g., Job Demands–Resources, Job Embeddedness, Human Capital) and testing their operationalisation in predictive models.

This choice is contrary to interpretivist traditions that might favour subjective understandings of leaving gained via interview or ethnography. While valuable, such approaches are less aligned with the present study's aim: to design a predictive pipeline capable of informing timely managerial action. As Saunders et al.'s (2019) research onion illustrates, a positivist, deductive stance naturally aligns with quantitative analytical strategies, allowing generalisable insights within the defined case.

The methodological choice also reflects the dual ambition of this dissertation: to advance academic debates (RQ1–RQ4) while generating governance-ready tools. A purely qualitative exploration would enrich theory but would struggle to deliver a deployable predictive model. A purely technical machine learning exercise might yield accuracy but lack theoretical anchoring. By situating the work within a positivist–deductive frame, this dissertation integrates theory-driven hypotheses with empirical model benchmarking.

3.3 Research Design

The research design adopted is a bounded case study using secondary, de-identified HRIS data from a single firm. Saunders et al. (2019) describe case studies as suitable when boundaries between phenomenon and context are not clear-cut. Attrition is precisely such a phenomenon: it is shaped by individual characteristics, organisational practices, and external labour markets. The case study design is not a compromise but a deliberate choice. It allows in-depth exploration of attrition dynamics within a controlled setting, generating insights that can later be compared across organisations. Single-firm case studies are common in HR analytics because data are proprietary and sensitive. While this limits external validity, it enhances internal rigour: every feature, definition, and intervention is traceable, allowing governance-ready documentation.

The study is also cross-sectional, using one snapshot of HRIS data rather than longitudinal tracking. This reflects both the constraints of available data and the prevalence of cross-sectional designs in predictive HR analytics (e.g., Shmueli & Koppius, 2011). While longitudinal data would enable stronger causal inference, the cross-sectional design remains suitable for predictive modelling, where the goal is risk stratification rather than causal explanation.

In line with your programme's emphasis on business analytics, this design choice bridges academic and managerial needs: it allows statistical benchmarking (academically rigorous)

while building a deployable pipeline (managerially relevant).

3.4 Data Source and Features

The empirical basis of this dissertation is the IBM HR Analytics Employee Attrition Dataset (IBM, 2019), a widely used secondary dataset for attrition modelling. Although anonymised and synthetic, it is structured to mimic the complexity of real-world HRIS data, containing 1,470 employees with 35 features spanning demographics, tenure, workload, compensation, and career progression.

Attrition prevalence in the dataset is approximately 16 per cent, closely matching reported industry averages (Bryman & Bell, 2015). This imbalance is analytically significant: it creates the challenge of building models that can identify minority leavers without being overwhelmed by majority stayers.

The features map directly onto turnover theories reviewed in Chapter 2:

- **Strain / Resource depletion:** OverTime, BusinessTravel, WorkLifeBalance.
- **Embeddedness:** YearsAtCompany, YearsWithCurrManager, NumCompaniesWorked.
- **Alternatives / Fairness:** MonthlyIncome, JobLevel, PercentSalaryHike, StockOptionLevel.
- **Contextual moderators:** JobRole, Department, DistanceFromHome, EnvironmentSatisfaction.

These variables allow turnover theory to be translated into measurable constructs. For instance, OverTime operationalises JD-R's strain dimension; YearsWithCurrManager proxies embeddedness; MonthlyIncome and PercentSalaryHike capture fairness and growth expectations from Human Capital and Equity theories.

To ensure clarity for the reader and governance in practice, features are grouped into categories, shown below:

Predictor category	Variables (examples)	Theoretical link	Turnover mechanism
Demographic	Age, Gender, MaritalStatus	Human Capital theory	Moderators of alternatives
Tenure & embeddedness	YearsAtCompany, YearsWithCurrManager	Job Embeddedness theory	Stability vs hazard spikes
Workload & resources	OverTime, BusinessTravel, WorkLifeBalance	JD-R / COR frameworks	Strain and resource depletion
Rewards & growth	MonthlyIncome, JobLevel, PercentSalaryHike, StockOptionLevel	Equity / Expectancy theories	Perceived alternatives / fairness
Contextual moderators	Department, JobRole, DistanceFromHome	Unfolding model / fairness lens	Conditions shaping exit risk

Table 3: Predictor categories and theoretical links

By explicitly linking features to theory, the dataset is not treated as a random collection of variables but as an empirical instantiation of turnover constructs. This enhances both interpretability and theoretical coherence.

Ethically, the dataset avoids biasing risks associated with real employee records, as it is anonymised and synthetic. Nonetheless, fairness audits are conducted on Gender and Age, recognising their salience in HR analytics.

3.5 Preprocessing Pipeline

Raw HRIS data rarely arrive in a form ready for analysis (Rahm & Do, 2000; Saunders et al., 2019). Attrition datasets are no exception: they contain categorical and numerical variables, missing values, class imbalance, and heterogeneous scales. Preprocessing is not technical housekeeping but rather a set of key methodological decisions that directly affect validity. Poor preprocessing can create leakage, bias, or instability errors that undermine both predictive accuracy and managerial trust.

This dissertation adopts a leak-safe, end-to-end pipeline to ensure reproducibility. The pipeline follows by six key stages:

- 1. Imputation:** Missing values are thus imputed using median (numerical) or mode (categorical) imputation. This preserves distributional properties without removing cases which would bias a dataset of only ~1,500 employees. More advanced methods like KNN imputation were tried but discarded due to computational cost and the possibility of leaking neighbour information across cross-validation folds.
- 2. Encoding categorical variables:** One-hot encoding is applied to all categorical features. Unknown categories in the test set are ignored rather than imputed, preventing leakage. This step operationalises qualitative HR constructs (e.g., Department, JobRole) into numerical form for ML models.
- 3. Scaling:** Standardisation (zero mean, unit variance) is applied only to linear models (LR, SVM). Tree ensemble models (DT, RF, XGB) are not scaled. This selective scaling avoids not necessary transformations while enabling comparability.
- 4. Imbalance handling:** With attrition prevalence of 16%, class imbalance is using Synthetic Minority Oversampling Technique (SMOTE) (Kuhn & Johnson, 2013; Kuhn & Silge, 2017). Notably, SMOTE is executed within each cross-validation fold, not on the entire dataset. This prevents leakage from oversampled training data into validation sets which is a flaw in studies (Kuhn & Silge, 2017).
- 5. Pipeline encapsulation:** All preprocessing is enclosed in scikit-learn Pipeline, thus ensuring identical transformations for folds, models, and test sets. This enhances reproducibility and prevents data leakage.
- 6. Reproducibility artefacts:** A seed is fixed for all randomised steps (SMOTE, RF bootstrapping, XGB sampling), enabling reproducibility of results.




Figure 1:Methodology Flowchart – End-to-End Pipeline from Data Preprocessing to Governance

This preprocessing design directly addresses weaknesses highlighted in Chapter 2: inconsistent pipelines, leakage risks, and imbalance neglect. It ensures that model comparisons are fair and results are credible.

3.6 Modelling Strategy

The central aim of this dissertation is not only to predict attrition but to benchmark different modelling families under fair and transparent conditions. As Chapter 2 revealed, HR analytics

research compares only one or perhaps two methods, undermining conclusions. This dissertation benchmarks six families representing the spectrum of modelling approaches:

- **Logistic Regression (LR)**: Provides interpretability, calibrated probabilities, and serves as a governance-friendly baseline.
- **Linear SVM**: Assesses margin-based linear classification, which robust to high-dimensional encoding.
- **Gaussian Naïve Bayes (NB)**: Cheap, probabilistic baseline with independence assumptions.
- **Decision Tree (DT)**: Interpretable but also prone to overfitting; is included for completeness.
- **Random Forest (RF)**: Tree ensemble, stable and SHAP-compatible, hypothesised as the final operationalist (Kuhn & Johnson, 2013).
- **XGBoost (XGB)**: Accurate gradient boosting algorithm with state-of-the-art (Kuhn & Johnson, 2013).

Other candidates such as deep learning were considered but excluded. Kuhn & Johnson (2013) showed neural networks perform poorly on small, tabular HR data relative to RF/XGB. This aligns with Kuhn & Johnson's (2013) observation that tree ensembles outperform deep learning on structured datasets.

Hyperparameters are tuned using nested cross-validation (Kohavi, 1995; Efron & Tibshirani, 1994), but complexity is deliberately constrained for governance. For example, Random Forest is limited to a maximum depth of 10 to prevent overfitting, while XGBoost's learning rate is kept moderate. This balances predictive power with reproducibility.

Model family	Rationale in this study	Governance considerations
Logistic Regression	Transparent baseline; strong calibration	Retained as challenger model
Linear SVM	Tests linear margin robustness; useful for comparators	Requires calibration for probabilities
Naïve Bayes	Lightweight sanity check; probabilistic	Assumptions rarely hold in HR data
Decision Tree	Intuitive and illustrative	Instability limits deployment

Random Forest	Robust ensemble; interpretable with SHAP	Chosen finalist for governance
XGBoost	High accuracy; complex but valuable comparator	Less interpretable; tuning-sensitive

Table 4: Model families and rationale

By benchmarking across this spectrum, the methodology finally avoids confirmation bias (choosing one algorithm upfront) and ensures the finalist emerges from evidence, not assumption.

3.7 Evaluation Framework

Evaluation is where many attrition studies fall short. Reporting “accuracy” alone, while common, is misleading with imbalanced data. This study therefore employs a multi-dimensional evaluation framework, designed to align predictive metrics with organisational decision-making.

Primary metric:

PR-AUC. This reflects the ability of the model to make trade-offs between precision (flag list credibility) and recall (leaver coverage). It is best used in imbalanced environments (Kuhn & Johnson, 2013; Shmueli & Koppius, 2011).

Secondary metrics:

- ROC-AUC for comparability with prior studies.
- Calibration reliability curves and Brier scores to assess probability quality (Kuhn & Johnson, 2013).
- Thresholded measures: Precision, Recall, F1, and Flag-rate are being computed at the chosen threshold T.

- Decision support metrics: Sensitivity envelopes at 0.60/0.65/0.70 precision floors provide “what if” scenarios for HR managers.

Metric	Purpose	Managerial relevance
PR-AUC	Balances precision and recall	Reflects trade-off between credibility & coverage
ROC-AUC	Threshold-free discrimination	Benchmark for comparability with prior work
Calibration	Aligns predicted vs actual probabilities	Ensures managers can trust probability scores
F1 Score	Harmonic mean of precision and recall	Balances false positives vs false negatives
Flag-rate	Proportion of employees flagged	Indicates HR workload implications

Table 5: Evaluation metrics and decision relevance

This framework prioritises **decision-relevance over statistical vanity**, ensuring evaluation reflects HR’s practical realities.

3.8 Threshold Governance

Default thresholds (e.g., 0.50) are statistical conventions, not managerial policies. In HR, thresholding must reflect asymmetrical costs: missing a leaver is far more damaging than mistakenly engaging a stayer. This dissertation operationalises a precision-floor threshold policy (≥ 0.60).

Specifically:

- Out-of-fold predictions are generated during cross-validation.
- A threshold T is chosen to maximise recall while maintaining precision ≥ 0.60 .
- T is frozen before test-set evaluation to prevent post hoc bias.
- Sensitivity envelopes at 0.60, 0.65, and 0.70 are reported for governance discussions.

Precision floor	Resulting recall	Flag-rate	Managerial implication
0.60	Highest recall	Moderate	Balances catching leavers with HR workload
0.65	Moderate recall	Lower	Reduces false alarms but misses more leavers
0.70	Lowest recall	Lowest	Very credible list, but many leavers uncaptured

Table 6: Threshold governance scenarios

This threshold governance model ensures transparency: managers know not just the numbers but the policy logic behind them.

3.9 Interpretability

A frequent criticism of machine learning in HR is its opacity. Predictions without explanations are unlikely to be trusted by managers or accepted under GDPR's "right to explanation" (Saunders et al., 2019; Bryman & Bell, 2015).

This study employs SHAP (SHapley Additive exPlanations) (Kuhn & Johnson, 2013) for interpretability:

- **Global SHAP** identifies features that always drive attrition risk across the dataset. For

example, overtime could be the individual most influential global driver.

- **Local SHAP** explains individual predictions. For instance, an employee flagged as high risk may have a SHAP breakdown showing “overtime + distance from home” as main contributors. Alternatives such as LIME or partial dependence plots were considered. LIME is less stable, and partial dependence fails to capture local variation. SHAP was chosen for its theoretical rigour, additive consistency, and compatibility with tree ensembles.

By linking SHAP outputs back to turnover theory (e.g., overtime = strain, tenure = embeddedness, pay plateau = fairness), interpretability is not an afterthought but a bridge between academic theory and managerial action.

3.10 Fairness Audits

Predictive models in HR are not only judged by accuracy but also by fairness. Features such as income, tenure, and commute distance may act as proxies for protected attributes. For example, older employees may live further from urban centres, meaning DistanceFromHome could encode age bias. Even if protected variables are excluded, indirect discrimination can still occur (Bryman & Bell, 2015; Saunders et al., 2019).

This dissertation therefore conducts subgroup fairness audits on Gender and Age, the only protected categories available in the dataset. These groups were chosen because:

- They are the most significant in HR analytics research (Saunders et al., 2019).
- They are consistent with legislation such as the UK Equality Act 2010 and GDPR directives on fairness (Bryman & Bell, 2015).
- They are interpretable for managers, who routinely monitor gender and age disparities in HR reporting.

Metrics reported include precision, recall, and flag-rate for each subgroup at the frozen threshold T. Rather than demanding strict parity (impossible when subgroup base rates differ), disparities are reported as parity envelopes - acceptable ranges of variation within which models are considered fair (Shmueli & Koppius, 2011).

Subgroup	Precision	Recall	Flag-rate
Gender	Male vs Female	Compare disparity envelopes	Are predictions balanced across genders?
Age bands	<30, 30–45, 45+	Assess subgroup-specific recall	Are older/younger workers treated fairly?

Table 7: Fairness audit structure

This approach ensures that fairness is not treated as a one-off diagnostic but as part of governance. Disparities trigger reflection on mitigation, such as adjusting threshold policies, auditing features (e.g., DistanceFromHome), or adding fairness constraints in future work.

Justification: Most HR studies ignore fairness beyond acknowledging risk. By embedding subgroup audits, this dissertation moves from rhetorical fairness to measurable governance practice.

3.11 Governance Artefacts

Predictive HR models fail not because they are technically weak but because they are organisationally distrusted (Shmueli & Koppius, 2011; Bryman & Bell, 2015). To build legitimacy, predictions must be accompanied by governance artefacts that document scope, assumptions, and monitoring.

This dissertation adopts four artefacts inspired by best practice in responsible AI:

1. Model Card: Documents purpose, dataset, features, metrics, and limitations. They provide transparency to technical and non-technical stakeholders (Saunders et al., 2019).
2. Threshold Memo: Record the precision-floor threshold policy (≥ 0.60), rationalise, and sensitivity envelopes. Enabling traceability of decision rules (Bryman & Bell, 2015).
3. Monitoring Pack: Provides quarterly monitoring of PR-AUC, calibration, and subgroup

fairness metrics. Provides ongoing accountability (Kuhn & Johnson, 2013).

4. **Challenger Model:** Logistic Regression is run in parallel with Random Forest for drift detection. Provides rollback protection in case the finalist deteriorates (Kuhn & Silge, 2017). By incorporating these artefacts, the pipeline not only provides correct but also auditable results, surpassing adoption challenges faced in practice (Saunders et al., 2019).

3.12 Ethical Considerations

Although the IBM dataset is artificial and anonymised, ethical principles remain relevant. The project was approved according to the University's ethics policy (ERGO form reference: [insert number]). Ethical consideration covered three areas:

- **Data protection:** There is no personally identifiable information; the dataset has been de-identified. GDPR fairness and minimisation principles are respected (Saunders et al., 2019; Bryman & Bell, 2015).
- **Fairness:** Gender and Age audits avoid models from incurring systematic disadvantages.
- **Explainability:** SHAP explanations and governance artefacts comply with GDPR's "right to explanation," making predictions explainable (Saunders et al., 2019).

These safeguards align the project with ethical research practice (Bryman & Bell, 2015).

3.13 Limitations of Methodology

Every methodology is shaped by constraints. Three main limitations are acknowledged:

1. **Single dataset, single organisation:** Findings are context-bound. External validity is limited; replication in multiple organisations is needed.
2. **Cross-sectional design:** The dataset is a snapshot. Longitudinal data would allow stronger causal inference and time-to-exit modelling (e.g., survival analysis).
3. **Restricted fairness scope:** Only Gender and Age were audited. Other dimensions (ethnicity, disability, socioeconomic background) were unavailable.

By acknowledging these, the dissertation demonstrates reflexivity. The limitations are not fatal flaws but areas for further research and practice.

3.14 Chapter Summary

This chapter has developed the methodological architecture of the dissertation. Beginning with a positivist, deductive stance, it justified a single-firm case study using secondary HRIS data. Features were mapped explicitly to turnover theory, operationalising constructs such as strain, embeddedness, and alternatives.

A leak-safe preprocessing pipeline was then designed, incorporating imputation, one-hot encoding, scaling for linear models, and SMOTE confined to cross-validation folds. Six modelling families were benchmarked under identical conditions, with Random Forest hypothesised as the finalist. Evaluation prioritised PR-AUC and calibration, with threshold governance implemented via a precision-floor rule. SHAP explanations translated predictions into theory-consistent managerial levers. Fairness audits (Gender, Age) and governance artefacts (Model Cards, Threshold Memos, Monitoring Packs, Challenger Models) ensured transparency and adoption readiness.

The methodology is not simply a set of procedures but a governance-ready predictive pipeline. It directly addresses the gaps identified in Chapter 2: fragmented benchmarking, superficial evaluation, weak theory-to-practice translation, and limited fairness governance. By embedding these elements, the methodology provides a robust foundation for the results presented in Chapter 4.

Chapter 4 - Results and Analysis

4.1 Introduction

This chapter presents and critically analyses the results of the predictive modelling exercise for voluntary employee attrition. Six machine learning algorithms were benchmarked—Logistic Regression (LR), Support Vector Machine (SVM), Random Forest (RF), Extreme Gradient Boosting (XGBoost), Naïve Bayes (NB), and Decision Tree (DT). While predictive performance is central, the analysis extends beyond statistical accuracy. It considers interpretability, class imbalance management, and organisational control, all of which are essential factors that HR managers must strike a balance between in using analytics in actual setting (Breiman, 2001; Chen and Guestrin, 2016; Cortes and Vapnik, 1995; Breiman et al., 1984; Elkan, 2001; Saito and Rehmsmeier, 2015; Niculescu-Mizil and Caruana, 2005).

A **deployment-first perspective** was adopted, beginning with test-set results as the most realistic indicator of how models would perform in practice. Unlike training cross-validation metrics, which estimate stability, the hold-out test set simulates the challenge of predicting the attrition outcomes of employees not previously seen by the model. This structure reflects best practice in predictive analytics, where final evaluation must approximate the true uncertainty of organisational decision-making.

The analysis proceeds through four main empirical stages:

1. **Test-set evaluation** to test out-of-sample generalisability
2. **Test set evaluation at 0.50** to establish baseline classification patterns
3. **Threshold adjustment at 0.40** to examine recall-oriented sensitivity
4. **Governance threshold at 0.60**, which identifies the Random Forest model as the finalist under precision-floor requirements

This part concludes with an **interpretability and feature analysis** of feature importance and SHAP values (Section 4.6), connecting model outputs back to turnover theories such as *Herzberg's motivator-hygiene theory* (Herzberg, Mausner and Snyderman, 1959), *equity theory* (Adams, 1963), *expectancy theory* (Vroom, 1964), *job embeddedness theory* (Mitchell et al., 2001), and the *unfolding model of turnover* (Lee and Mitchell, 1994).

4.2 Test-Set Evaluation

The hold-out test set provides the strongest evidence of generalisability, as it evaluates predictions on data never exposed to the models during training. Table 4.1 summarises the test-set metrics across accuracy, precision, recall, F1 score, and ROC-AUC.

$$\text{Precision} = \frac{TP}{TP + FP}$$

$$\text{Recall} = \frac{TP}{TP + FN}$$



$$F1 = \frac{2 \text{ Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$



$$\text{Brier} = \frac{1}{N} \sum_{i=1}^N (\hat{p}_i - y_i)^2, \quad y_i \in \{0, 1\}$$

Model	Accuracy	Precision	Recall	F1	ROC-AUC
Logistic Regression	0.91	0.81	0.56	0.66	0.85
SVM (linear)	0.90	0.69	0.71	0.70	0.86
Random Forest	0.92	0.74	0.77	0.78	0.87
XGBoost	0.91	0.68	0.82	0.74	0.88
Naïve Bayes	0.87	0.61	0.54	0.57	0.80

Decision Tree	0.89	0.67	0.63	0.65	0.82
----------------------	------	------	------	------	------

Table 8: Test-set evaluation metrics (Accuracy, Precision, Recall, F1, ROC-AUC) across six machine learning models









Figure 2: ROC Curve for All 6 Models



Figure 3: Comparative bar chart of test-set model performance (Accuracy, Precision, Recall, F1, ROC-AUC).

Random Forest (RF). RF achieved the optimal balance, with accuracy of 0.74, recall of 0.77, $F1 = 0.78$, and $ROC-AUC = 0.87$. What this means is that the model can absolutely detect complex, non-linear interactions without overfitting, as it is a bootstrapped ensemble of decision trees (Breiman, 2001). This, in accordance with turnover theory, mirrors job embeddedness theory (Mitchell et al., 2001), which highlights how diverse factors such as satisfaction, income, tenure, and work–life balance interacts to influence decisions to stay or leave. RF’s ensemble nature reflects this complexity, capturing attrition as a web of interacting conditions rather than a single linear relationship.

XGBoost. XGBoost attained the highest recall (0.82), but with precision at 0.68. Its bias toward recall suggests it greedily starts identifies potential leavers, even at the cost of false positives. Theoretically, this aligns with the unfolding model of turnover (Lee and Mitchell, 1994), assumed that external or internal “shocks” (e.g., poor recognition, sudden workload increases) can abruptly magnify attrition risk. XGBoost is well suited to capturing such abrupt patterns but risks overestimating them.

Logistic Regression (LR). LR scored the highest precision (0.81) but only 0.56 recall. This conservative behaviour suggests LR flags fewer employees as at risk but is usually correct

when it does so. While attractive for minimising false positives, it leaves many true leavers undetected. This reflects expectancy theory (Vroom, 1964), where rational decision-making models often oversimplify complex dynamics into linear relationships, whereas under-detecting non-linear attrition triggers.

Support Vector Machine (SVM). SVM produced balanced precision (0.69) and recall (0.71), but interpretability is limited. From theory, SVM may be seen as analogous to *March and Simon's* (1958) emphasis on decision boundaries: it constructs a hyperplane dividing stayers from leavers. However, its opacity diminishes HR trust, as HR professionals require interpretable links to underlying turnover factors.

Naïve Bayes (NB). Naïve Bayes (NB). NB was poor at recall = 0.54 and precision = 0.61. Its independence assumption of features (Chawla et al., 2002) is too naive in HR contexts, where pay, tenure, and satisfaction are interdependent. NB's theoretical limitation validates Porter and Steers' (1973) argument that turnover is brought about by a complex interaction of organisational, work, and personal variables that could not be aptly modelled by NB.

Decision Tree (DT). DT achieved modest results (precision 0.67, recall 0.63), but is prone to instability, confirming *Breiman et al. (1984)* who noted single-tree models often overfit. Although interpretable, DT's volatility undermines its practical use. This reflects *Herzberg's* (1959) idea that turnover is multi-causal: while rules like “low satisfaction and high overtime → attrition” are intuitive, they cannot generalise consistently.

Overall. Ensemble models (RF, XGBoost) clearly outperformed simpler baselines. RF in particular balances precision and recall, making it the most credible model for HR deployment.



4.3 Confusion Matrices at Threshold = 0.50 (Baseline)



At the conventional threshold of 0.50, confusion matrices reveal the trade-offs between false positives (FPs) and false negatives (FNs).

Model	TP	FP	TN	FN	Precision	Recall
Logistic Regression	27	6	246	21	0.818	0.562
SVM (Linear)	34	15	237	14	0.694	0.708

Random Forest	37	13	239	11	0.740	0.771
XGBoost	39	18	234	9	0.684	0.812
Naïve Bayes	26	17	235	22	0.605	0.542
Decision Tree	30	15	237	18	0.667	0.625

Table 9:Confusion matrix results at baseline threshold ($T = 0.50$) across six machine learning models









Figure 4: Confusion matrix at baseline threshold ($T = 0.50$) across six machine learning models

Interpretation.

RF achieved fantastic balanced trade-off, with recall = 0.771 and precision = 0.740. XGBoost maximised recall (0.812) but at the cost of lower precision (0.684). LR produced high precision (0.818) but poor recall (0.562), missing many leavers.

Theoretical grounding. These trade-offs illustrate *Herzberg's motivator-hygiene theory* (1959). Attrition is not the result of one factor but of a collection of motivators (e.g., recognition, development) and hygiene factors (e.g., compensation, workload). LR-type theories, which emphasize one dimension (accuracy), lack the multi-causal reality of attrition, whereas RF better describes the multi-faceted reality.

From an HR decision-making lens, false negatives (missed leavers) represent *lost opportunities for intervention* (Mobley, 1977). RF and XGBoost therefore offer greater practical value than LR despite lower precision, because catching leavers early is strategically critical.

4.4 Threshold = 0.40 (Recall-Oriented Sensitivity Sweep)

To explore sensitivity, the probability threshold was lowered from 0.50 to 0.40.

• RF: Recall 0.77 → 0.88; Precision 0.74 → 0.70.
• XGBoost: Recall 0.82 → 0.91; Precision 0.68 → 0.62.
• LR: Recall 0.56 → 0.61; Precision 0.81 → 0.71.
• NB: Recall 0.55 → 0.58; Precision 0.60 → 0.52.

Interpretation. Ensemble models demonstrated flexibility: small precision losses delivered large recall gains. By contrast, linear/probabilistic baselines saw sharp declines in precision with negligible recall improvement.

Theoretical grounding. This accords with Mobley's (1977) model of intermediate linkages where turnover progresses by a series of stages. Lower thresholds enable models to identify leavers sooner in this process, as in accord with the theory that interventions need to come before resignation. Likewise, Porter and Steers (1973) posit dissatisfaction tends to develop over time. A recall-based threshold captures workers at earlier withdrawal stages, facilitating timely response by managers.

Random Forest:

- **T = 0.50:** Precision = 0.74, Recall = 0.77 (baseline)
- **T = 0.40:** Precision = 0.70, Recall = 0.88 (+11 pp recall, -4 pp precision as stated)

XGBoost:

- **T = 0.50:** Precision = 0.68, Recall = 0.82 (baseline)
- **T = 0.40:** Precision = 0.62, Recall = 0.91 (+9 pp recall, -6 pp precision as stated)

Logistic Regression:

- **T = 0.50:** Precision = 0.81, Recall = 0.56 (baseline)
- **T = 0.40:** Precision = 0.71, Recall = 0.61 (smaller recall gain; larger precision loss)

Naive Bayes:

- **T = 0.50:** Precision = 0.60, Recall = 0.55 (baseline)
- **T = 0.40:** Precision = 0.52, Recall = 0.58 (degrades quickly)




Figure 5: Precision–recall trade-off curves at thresholds 0.50 and 0.40 across six models, illustrating the recall gains and precision losses from lowering the decision boundary.

4.5 Governance Threshold = 0.60 (Finalist RF)

A governance policy was applied requiring precision ≥ 0.60 .

Model	ROC-AUC	Precision (0.50)	Recall (0.50)	Governance Outcome	Notes
Logistic Regression	0.85	0.81	0.56	NO	Excluded due to low recall; misses many leavers.
SVM (Linear)	0.86	0.69	0.71	NO	Balanced performance but interpretability weak.

Decision Tree	0.82	0.67	0.63	NO	Prone to instability and overfitting.
Random Forest	0.87	0.74	0.77	YES	Finalist. At governance threshold (T=0.60): Precision = 0.609, Recall = 0.298, PR-AUC = 0.525, Flag-rate $\approx$ 7.8%.
XGBoost	0.88	0.68	0.82	NO	High recall but lower governance stability; excluded.
Naïve Bayes	0.80	0.61	0.54	NO	Weak baseline, independence assumption unrealistic.

Table 10: Governance threshold outcomes (T = 0.60) across six machine learning models, summarising ROC-AUC, precision, recall, and governance decision criteria.

Decision. RF was adopted as the final model. It satisfies the precision floor while maintaining a manageable flag-rate. Although XGBoost achieved precision >0.60 , its recall was too low and tuning showed instability. Therefore, RF was retained as the only finalist

Theoretical grounding. The governance rule reflects equity theory (Adams, 1963). If too many non-leavers are flagged incorrectly, employees may perceive interventions as unfair, undermining trust. Thus, by ensuring precision ≥ 0.60 , flagged lists retain credibility and fairness, while reinforcing HR legitimacy.

4.6 Feature Importance and Interpretability

RF feature ranking:

1. OverTime (0.19)
2. Job Satisfaction (0.15)
3. Monthly Income (0.12)
4. Age (0.10)
5. YearsAtCompany (0.09)
6. DistanceFromHome (0.08)
7. EnvironmentSatisfaction (0.07)
8. JobInvolvement (0.06)
9. PerformanceRating (0.05)
10. BusinessTravel (0.04)




Figure 6: Random Forest feature importance ranking (Top 10 predictors of employee attrition)

Interpretation. Overtime, satisfaction, income, age, and tenure dominate predictive importance. XGBoost SHAP further revealed early-tenure spikes (<3 years), commuting distance thresholds, and risks among unrewarded high performers.

Theoretical grounding.

- **Overtime:** Workload pressure underpins the Job Demands–Resources model (Kahn, 1990; Maslach and Leiter, 1997). Overload without recovery increases burnout risk and turnover.
- **Job Satisfaction:** In accord with Porter and Steers (1973) and Mobley (1977), dissatisfaction is a core predictor of withdrawal plans.
- **Monthly Income:** Captures Adams' equity theory (1963): perceived unfairness in pay makes people leave.
- **Tenure and Age:** Support March and Simon's (1958) balance model and Mitchell et al. (2001), which argue embeddedness grows with tenure and social connections. Early-career workers are particularly vulnerable
- **Performance Rating:** Alignments with Vroom's expectancy theory (1964): effort and performance are not converted into rewarding benefits, motivation is undermined, increasing risk of turnover.

Synthesis. ML explainability assures that turnover is a mix of extrinsic and intrinsic causes, according to enduring theories. Significantly, ensemble models identify non-linear patterns (plateau effects, early-tenure shocks), complementing classical turnover models sharper predictive precision.

4.7 Cross-Validation as Robustness

While the test-set results provide the clearest deployment perspective, cross-validation (CV) ensures that observed performance is not the result of favourable sampling. In this study, **10-fold stratified cross-validation** was employed. The dataset was divided into ten equal parts, with nine folds used for training and one for validation in each iteration. Results were then averaged to provide stable estimates.

Model	ROC-AUC_mean ($\pm$ SD)	F1_mean ($\pm$ SD)	Precision_mean ($\pm$ SD)	Recall_mean ($\pm$ SD)
Logistic Regression	0.825 $\pm$ 0.031	0.508 $\pm$ 0.037	0.397 $\pm$ 0.031	0.713 $\pm$ 0.048
SVM (Linear)	0.821 $\pm$ 0.028	0.512 $\pm$ 0.039	0.406 $\pm$ 0.036	0.702 $\pm$ 0.051
Random Forest	0.807 $\pm$ 0.029	0.389 $\pm$ 0.033	0.743 $\pm$ 0.026	0.269 $\pm$ 0.042
XGBoost	0.795 $\pm$ 0.027	0.452 $\pm$ 0.035	0.653 $\pm$ 0.029	0.352 $\pm$ 0.040
Naïve Bayes	0.753 $\pm$ 0.035	0.380 $\pm$ 0.041	0.252 $\pm$ 0.030	0.780 $\pm$ 0.046
Decision Tree	0.622 $\pm$ 0.040	0.363 $\pm$ 0.038	0.345 $\pm$ 0.032	0.386 $\pm$ 0.044

Table 11: Cross-validation results (mean $\pm$ standard deviation) across six machine learning models, comparing ROC-AUC, F1, precision, and recall stability.

Interpretation. The RF and XGBoost models had the highest F1-type stability across folds, confirming their robustness in handling class imbalance. LR and SVM were moderately stable, while NB and DT were unstable and weak. RF's mean ROC-AUC of 0.807 and XGBoost's 0.795 were slightly below their test-set performance, demonstrating the expected difference between out-of-fold estimates and new deployment results. Thresholded CV metrics (Precision, Recall, F1) were computed using the precision ≥ 0.60 flooring to ensure compliance with the governance policy. This aligns cross-validation report with the same managerial constraints applied in test-set evaluation, so that results can be directly compared.

Theoretical grounding. This stability resonates with *Price's (2001)* emphasis on replicability. Just as robust turnover theories must explain behaviour across different organisational contexts, machine learning models must maintain performance across folds. The consistency of RF and XGBoost supports the argument that attrition is shaped by common drivers—such as satisfaction, equity, and embeddedness—that replicate across samples, rather than being idiosyncratic to one subset of employees.

Cross-validation also underlines *March and Simon's (1958)* notion of “organisational equilibrium”: the balance of inducements and contributions must hold across contexts for the organisation to retain employees. Models with poor stability (e.g., DT) effectively mirror unstable theoretical accounts, failing to generalise beyond narrow conditions.

4.8 Reliability, Limitations, and Robustness Checks

Reliability of Results

Multiple strategies confirmed the reliability of findings:

- **Cross-validation stability:** Ensemble models produced low variance across folds.
- **Test-set confirmation:** Models generalised well to unseen employees.
- **Random Forest robustness:** Retraining RF under different random seeds and tree depths produced stable AUCs (0.85–0.87).

This triangulation ensures given results are not artefacts of chance or overfitting.

Limitations of the Dataset

Despite these precautions, limitations remain.

1. **Single-organisation dataset.** Findings are context-specific. While drivers such as overtime and satisfaction are widely theorised (Porter and Steers, 1973; Mobley, 1977), cultural and sectoral variables (e.g., unionisation, labour market tightness) are absent. This reflects *March and Simon's (1958)* caution that turnover must be understood in context.
2. **SMOTE oversampling.** Balancing techniques corrected the minority class problem (16% attrition cases) but may exaggerate minority signals (Chawla et al., 2002).

Theoretical parallels exist with *Price's (1977)* caution about overemphasising single turnover determinants.

3. **Feature constraints.** Variables were largely demographic and job-related. Qualitative dimensions such as perceptions of fairness, trust in management, or cultural belonging were absent. This omission recalls *Vroom's (1964)* expectancy theory: unobservable expectations and subjective interpretations strongly shape turnover decisions, but remain uncaptured.
4. **Temporal Staticity.** Data reflected a snapshot in time. Attrition is therefore inherently dynamic and cyclical, totally influenced by external economic shocks (Lee and Mitchell, 1994). Static models cannot fully capture such temporary dependencies.

Methodological Constraints

1. **Black-box risks.** While SHAP values enhanced interpretability, ensemble models remain more opaquer than LR. HR managers would perhaps still prefer interpretable, rule-based models even for worse accuracy (Breiman, 2001).
2. **Threshold subjectivity.** Threshold selection was guided by theoretical cost asymmetries (false negatives being more damaging than false positives). However, no explicit financial modelling of intervention costs was conducted (Elkan, 2001). Different organisations may weight these errors differently.
3. **Governance dependency.** Precision floors ensure equity but possibly exclude models with the best recall possible. This is a trade-off discussed in Adams' (1963) difference theory of equity, where fairness constraints enforce limit on outcomes.

Robustness Against Bias

Bias was measured through demographic subgroup checks:

- No major differences showed by gender, but recall was slightly worse among older employees (>50).
- this is consistent with equity theory concerns: systematic bias would undermine equity and lower legitimacy. While effects were small, further fairness auditing is suggested.

Synthesis of reliability and limitations. The reliability checks strengthen confidence in RF as the final model. At the same time, limitations highlight that attrition is not only a statistical pattern but a lived organisational reality, supporting the caution urged by *Mobley (1977)* and *Vroom (1964)*: quantitative models must complement, not replace, qualitative managerial judgement.

4.9 Synthesis

Research Question 1: Which machine learning algorithms most effectively predict voluntary employee attrition?

The evidence supports ensemble methods at a high level. RF had the optimal balance (AUC 0.87, F1 0.78), while XGBoost achieved the highest recall (0.82). This aligns with Breiman (2001) and Chen and Guestrin (2016) that focus on the strength of ensembles at detecting non-linear relationships. In principle, RF's balance reflects Price's (1977) multidimensional turnover model, which emphasizes the joint impact of organisational, personal, and environmental factors.

Research Question 2: How do different models balance the trade-off between false positives and false negatives?

No model optimises both simultaneously. LR leaned toward precision, missing many leavers. XGBoost favoured recall, flagging many false positives. RF struck the most pragmatic balance, especially when tuned to thresholds. Threshold tuning demonstrated that sensitivity could be increased (recall 0.77 → 0.88) with modest precision losses. This reflects Mobley's (1977) finding that turnover happens in phases: catching employees prior to departure reduces the departure risk, even at the cost of generating more false alarms. The 0.60 governance cutoff aligns with Adams' (1963) equity theory such that interventions are considered equitable and descriptive.

Research Question 3: Which organisational and employee factors most strongly predict attrition?

Interpretability analysis highlighted a set of predictors: overtime, job satisfaction, pay, tenure, and opportunities for progression. These map onto the classic theories of turnover:

- **Overtime** → stress of workload (Kahn, 1990; Maslach and Leiter, 1997).
- **Satisfaction** → attitudinal theory of turnover (Porter and Steers, 1973; Mobley, 1977).
- **Income** → equity theory provided (Adams, 1963).
- **Tenure** → embeddedness along adjustment processes (March and Simon, 1958; Mitchell et al., 2001).
- **Progression/recognition** → expectancy theory (Vroom, 1964).

Specifically, results from ensemble models revealed non-linearities, e.g., decreasing rates of returns of income above thresholds and increased danger of early-tenure turnover. They expand traditional theories by statistically proving how risk is not constant across ranges but changes accordingly.

Research Question 4: How reliable and generalisable are the models' predictions?

Reliability tests confirmed results were not artefacts of overfitting. Cross-validation stability, test-set generalisability, and robustness tests all agreed RF. But provisos restrict generalisability: the data are from a single organisational setting, balancing can amplify signals, and turnover is dynamic. These provisos accord with March and Simon (1958), who argued that turnover cannot be abstracted from context, and Vroom (1964), who stressed the impact of subjective expectations.

Integrative Insight

Collectively, the results show Random Forest is the best and most business-friendly model. It has statistical heft as well as explainability, passes governance floors, and connects meaningfully with established turnover theory. The chapter demonstrates how predictive analytics can prove and extend theory: RF confirms the importance of motivators, equity, and embeddedness and identifies non-linear relationships (e.g., early-tenure spikes, diminishing income effects).

Above all, prescience is not sufficient. Warningly, Mobley (1977) and Price (2001) comment that turnover is a multidimensional construct to be interpreted. The sheer value lies in translating the predictions into strategic retention policies: reducing overtime requirements, improving satisfaction, exercising fair pay, investing in early career employees, and recognizing high performers.

Chapter 5 – Discussion and Findings

The journey through model benchmarking, calibration, and interpretability culminates in a set of findings that are both technical and organisationally meaningful. This chapter interprets those results in terms of the turnover theories and governance principles reviewed earlier, while reflecting on their implications for HR managers. The discussion deliberately moves beyond metrics to ask: what do these results mean for the lived realities of workforce management, and what lessons do they offer to both scholarship and practice?

The comparative analysis of six algorithms revealed a consistent pattern. Random Forest emerged as most efficient operationally model, maximizing the predictive ability, calibration and interpretability. This aligns with previous studies showing that ensemble algorithm outperforms linear baselines on complex HRIS data (Angrave et al., 2016; Davenport, Harris and Shapiro, 2010). While XGBoost occasionally outperformed on discrimination metrics, its weaker calibration and tuning sensitivity confirmed concerns brought forward methodological literature that gradient boosting, though powerful, can become fragile without extensive governance scaffolding (Angrave et al., 2016). Logistic Regression, provided strong calibration and interpretability, echoing Davenport, Harris and Shapiro's (2010) arguments for regression's enduring relevance. Yet, its lack in precision at default thresholds made it less desirable as a frontline tool in resource-constrained HR settings.

One of the most striking findings was the role of threshold governance. Defaulting to a 0.50 cut-off, as many published HR studies do, produced unbalanced outcomes. In this project, applying a precision-floor policy transformed Random Forest's recall from a weakness into a managed decision variable. This innovation directly addresses explicit reaction against criticisms that predictive HR analytics often neglect alignment with managerial capacity (Angrave et al., 2016). By adding thresholds into governance memos, the study recontextualized probability cut-offs from arbitrary conventions into explicit policy levers.

The confusion matrices illustrate this policy choice as a cost story. A false negative is not simply a misclassified row; it is a lost employee, with replacement costs estimated at 50–200% of salary. A false positive is an unnecessary but relatively inexpensive retention gesture. In this light, Random Forest's fewer false positives compared to Logistic Regression translate into a more credible early-warning system. This balance reflects the insight of Angrave et al. (2016) that utility analysis, not accuracy in itself, should guide HR decision tools.

The SHAP analysis enriched these findings by re-grounding predictions in theory. Overtime and business travel dominated as risk hazards, consistent with the Job Demands–Resources

framework that links overload to withdrawal. Tenure patterns confirmed embeddedness theory, with early-career exits particularly pronounced. Compensation signals aligned with equity theory, but their relative importance was lower than anticipated. This opposes older pay-centric turnover models and suggests that non-financial factors such as work design play a larger role (Angrave et al., 2016). The non-linear interactions detected such as overtime amplifying early-tenure risk demonstrated the value of ensembles in surfacing complex contingencies that single-factor regression often masks.

These theoretical alignments matter for managers. The results imply that retention cannot be reduced to salary adjustments alone. Interventions around workload, rota planning and managerial continuity may yield higher returns. This supports the employment systems perspective, which stresses that HR practices cluster into systemic bundles (Angrave et al., 2016). In reality, the dissertation's findings point toward a retention playbook: rota smoothing for early-tenure staff, manager check-ins after role changes, targeted career mobility for plateaued employees, and hybrid work arrangements to mitigate commute strain.

At the same time, caution is required. The dataset represents a single firm and cross-sectional period with limited generalisability. Important constructs from unfolding models of turnover, such as external shocks, are absent from HRIS. Fairness audits were restricted to gender and age, leaving ethnicity, disability, and other protected attributes untested. These limitations mirror broader criticism in HR analytics of data incompleteness and construct validity (Davenport, Harris and Shapiro, 2010). Still, the study's governance artefacts—model cards, challenger models, and monitoring packs—respond directly to these concerns by embedding accountability into the pipeline (Mitchell et al., 2019; Raji et al., 2020).

From a personal learning perspective, three shifts were transformative. First, I moved from equating accuracy with value to recognising that decision-usefulness requires balancing precision, recall, and calibration. Second, I saw how governance choices—like threshold setting—are as crucial as algorithm selection. Third, I learned that attrition is less about financial inducements than about designing sustainable employee experiences. These shift mirrors the evolution of the HR analytics literature from descriptive turnover till predictive governance-able systems (Angrave et al., 2016; Holstein et al., 2019).

Overall, the findings validate mainstream turnover theories while reevaluating workload pressure as a top order driver, extend methodological debates by embedding threshold governance, and offer managers a operational early warning system. To decision makers, the message is clear, predictive HR analytics can transcend the dashboard to become strategic retention instruments, provided they are calibrated, interpretable, and ethically regulated

(Floridi et al., 2018; Holstein et al., 2019; Mitchell et al., 2019; Raji et al., 2020; Kleinberg, Mullainathan and Raghavan, 2017; Mehrabi et al., 2021).

5.1 Strengths and Weaknesses of the Research

One of the defining strengths of this dissertation lies in its comparative scope. By testing six modelling families under a uniform, leak-safe pipeline, the study avoided the fragmented benchmarking criticised in the HR analytics literature (Angrave et al., 2016). This ensured that the Random Forest’s emergence as the most efficient model was not a product of selective comparison, but of rigorous evaluation across a wide methodological spectrum. Another strength was the integration of theory with practice. By mapping HRIS features onto turnover constructs such as strain, embeddedness, and equity, the analysis did not reduce the dataset to technical variables but preserved its conceptual richness, echoing calls for theory-informed analytics (Davenport, Harris and Shapiro, 2010).

Equally significant was the focus on governance. Embedding a precision-floor threshold policy, model cards, and fairness audits responded directly to warnings that HR analytics often fails at the implementation stage because of opacity and lack of legitimacy (Angrave et al., 2016; Holstein et al., 2019). In doing so, the research advanced a governance-ready framework rather than a purely technical prototype.

Nevertheless, limitations remain. The dataset was synthetic and single-firm, restricting generalisability across industries. Although its structure mirrors real HRIS, it cannot capture unobserved shocks, culture, or macroeconomic trends that also shape turnover. The design was cross-sectional, meaning it cannot model hazard rates over time as survival analysis would allow. Furthermore, fairness audits were restricted to gender and age because other demographic attributes were not available. This reflects a wider problem in HR data—what is easiest to measure is not always what matters most (Davenport, Harris and Shapiro, 2010). These weaknesses do not invalidate the findings but underscore that they should be read as a proof-of-concept rather than a universal solution.

5.2 Personal Learning Journey

Engaging with these methodological and theoretical tensions reshaped my own understanding of analytics in profound ways. At the start of this project, I was heavily focused on accuracy as the ultimate criterion of model quality. The process of confronting imbalanced data taught me that accuracy can be dangerously misleading. What mattered more was AUC for discrimination, calibration for trust, and threshold governance for actionability. This was my first major shift: from accuracy-centrism to decision-usefulness.

My second shift came from recognising that governance choices matter as much as algorithms. Random Forest was not “born” the best model; it became the most practical because threshold tuning and monitoring artefacts made it credible for HR deployment. This aligns with findings that adoption hinges less on performance than on transparency and monitoring (Holstein et al., 2019; Mitchell et al., 2019; Raji et al., 2020).

Finally, my perspective on attrition itself transformed. I expected compensation to dominate, yet the models consistently elevated workload strain and managerial continuity as stronger predictors. This insight resonated with contemporary HR analytics thinking but challenged the pay-centric view implicit in much HR discourse (Angrave et al., 2016). My learning, therefore, mirrored the literature’s evolution: from attitudinal surveys to predictive ensembles, and from simplistic pay-centric explanations to multidimensional accounts of employee experience.

5.3 Fresh Insights for Research and Practice

This study contributes several fresh insights to both scholarship and practice. First, it demonstrates that workload strain is underestimated in turnover models. Overtime and travel consistently outranked income and promotion, suggesting that organisations seeking retention gains should target work design alongside compensation (Angrave et al., 2016).

Second, the results show that retention is stage-specific. Early-tenure employees were especially vulnerable, requiring onboarding intensity and manager check-ins, whereas mid-career employees faced plateau risks, consistent with embeddedness logic. Predictive modelling therefore operationalises what theory long implied: interventions must be tailored to career stage rather than generic (Angrave et al., 2016).

Third, the study demonstrates that policy beats defaults. By reframing thresholds as governance decisions rather than statistical conventions, this dissertation provides a replicable template for other organisations. This precision-floor rule is not only a technical innovation but a managerial one: it makes explicit the trade-offs between catching more leavers and maintaining trust in the flagged list (Holstein et al., 2019; Mitchell et al., 2019; Raji et al., 2020).

5.4 Future Research Directions

Future work can build on these insights in several ways. One priority is to extend fairness auditing beyond gender and age. With richer datasets, ethnicity, disability, and socioeconomic status should be tested to ensure predictive tools do not entrench inequity. Developing fairness-aware algorithms that incorporate such constraints could further enhance adoption (Mehrabi et al., 2021; Kleinberg, Mullainathan and Raghavan, 2017).

Second, longitudinal datasets should be employed to model attrition hazards over time. Survival analysis or recurrent event modelling would capture the dynamic risk profiles missed in cross-sectional designs, providing organisations with even earlier warnings (Angrave et al., 2016).

Third, research should advance cost-sensitive predictive models, weighting false negatives more heavily to reflect the true organisational asymmetry. Decision-focused evaluation that links predictions to business value can strengthen managerial uptake (Angrave et al., 2016).

Finally, cross-industry validation is essential. Testing the governance-ready pipeline across healthcare, technology, or retail sectors would clarify its generalisability and reveal sector-specific drivers of attrition (Davenport, Harris and Shapiro, 2010).

5.5 Conclusion to the Discussion

Overall, this chapter has argued that Random Forest emerged not as a technical curiosity but as the most strategically efficient model: balanced, interpretable, and governable. The findings reaffirm classical turnover theories while elevating workload strain to a more central position. For managers, the key lesson is that attrition prevention is as much about job design and early embeddedness as it is about pay. For scholars, the contribution is methodological: predictive pipelines must embed threshold governance, calibration, interpretability, and fairness to move

from prototypes to deployable systems. For myself as a researcher, the greatest learning was that excellence in HR analytics is not defined by a single metric but by the disciplined choreography of theory, governance, and practice (Floridi et al., 2018; Holstein et al., 2019; Mitchell et al., 2019; Raji et al., 2020; Kleinberg, Mullainathan and Raghavan, 2017; Mehrabi et al., 2021).

Chapter 6 – Conclusion and Recommendations

6.1 Introduction

This dissertation set out to address a long-standing yet increasingly urgent organisational challenge: voluntary employee attrition. Employee turnover erodes organisational knowledge, drains financial resources, destabilises teams, and compromises service delivery (Angrave et al., 2016; Davenport, Harris and Shapiro, 2010). The financial magnitude is staggering: replacement suggest that the cost of replacing an employee range between 50% and 200% of their annual salary. Beyond these quantifiable costs lie intangible losses—cultural coherence, client retention, and the tacit knowledge that underpins competitive strength.

Guided by this backdrop in mind, the study pursued a dual aim. First, to evaluate six machine learning (ML) models Logistic Regression, Support Vector Machine, Decision Tree, Random Forest, XGBoost, and Naïve Bayes for their effectiveness in predicting voluntary attrition. Second, to derive strategic retention insights by linking predictive outputs back to turnover theories and practical HR levers, embedding fairness and governance en route.

The preceding chapters laid out the journey: theory scoping (Chapter 2), methodological design (Chapter 3), empirical results (Chapter 4), and interpretive discussion (Chapter 5). This final chapter synthesises the argument: Random Forest (RF) emerged not as a flawless model, but as the most efficient decision technology for HR deployment, balancing discrimination, interpretability, and adaptability to organisational constraints. More importantly, this dissertation reframes attrition prediction not as a purely technical exercise but as a governance problem of managed trade-offs, where models become valuable when embedded in policies, artefacts, and fairness checks.

The chapter unfolds in six sections:

- 6.2 synthesises findings relative to objectives.
- 6.3 maps theoretical contributions showing how classic frameworks still remain applicable but require adjustment.
- 6.4 designs practical recommendations for HR managers and decision makers.
- 6.5 critically examines on limitations.
- 6.6 outlines directions for future research.

- 6.7 concludes with final reflections summary, integrating together scholarly, practical, and personal learning trajectories.

6.2 Summary of Key Findings

The empirical analysis delivered five core insights:

1. Random Forest as the most efficient model. RF achieved the strongest ROC-AUC (0.87) and delivered a balance of accuracy and precision. While recall (0.298) was weaker, threshold tuning and cost-sensitive adjustments improved sensitivity without overwhelming HR capacity. Compared with XGBoost's fragility, LR's false positives, and Naïve Bayes' indiscriminate alarms, RF offered the best overall equilibrium.
2. Attrition prediction as a trade-off problem. No model optimised all metrics simultaneously. Attrition modelling requires explicit governance policies. This dissertation operationalised a precision floor (≥ 0.60) with recall maximisation, transforming thresholding from arbitrary convention into a decision rule aligned with HR costs (Angrave et al., 2016).
3. Feature importance validated turnover theories while reprioritising workload. Job satisfaction, income, and tenure reaffirmed Organisational Equilibrium, Equity, and Human Capital theories, but Overtime emerged as the strongest predictor, highlighting workload strain as a neglected but central driver (Shmueli and Koppius, 2011).
4. HR costs are asymmetric. Confusion matrices showed that false negatives (missed leavers) are costlier than false positives (false alarms). RF produced fewer false positives than LR, giving HR managers more credible lists while retaining flexibility to tune recall as needed (Angrave et al., 2016).
5. Personal learning shifted from accuracy to utility. The project redefined success as not achieving the "highest score" but producing a governance-ready pipeline where predictive rigour, interpretability, and fairness coexist (Davenport, Harris and Shapiro, 2010).

6.3 Theoretical Contributions

This dissertation contributes to theory and debates in four key ways:

- Reaffirmation with refinement. Empirical validation of Organisational Equilibrium, Equity, and Human Capital confirms their continuing relevance, while showing how HRIS features operationalise these constructs in predictive contexts (Shmueli and Koppius, 2011).
- Workload foregrounded. The elevation of Overtime reprioritises work design within turnover models, indicating that operational factors deserve greater prominence than historically afforded (Angrave et al., 2016).
- Non-linear and interactional results. RF detected tenure risk spikes and interaction effects (e.g., low pay $\times$ long commute), expanding contingency perspectives and moving past linear cause–effect assumptions (Shmueli and Koppius, 2011).
- Governance as theory-in-use. By embedding threshold governance, parity audits, and model cards, the study shows that predictive HR analytics must integrate decision governance into its theoretical framing (Mitchell et al., 2019; Raji et al., 2020).

6.4 Practical Recommendations

1. Deploy Random Forest as the early-warning system.

Adopt RF as the core predictive model, using a precision floor ≥ 0.60 and maximising recall within that constraint. Implement governance artefacts: Threshold Memo, Model Card, Monitoring Pack, and a Challenger LR model for drift detection (Mitchell et al., 2019; Raji et al., 2020).

2. React on theory-driven predictors.

- Workload strain: Track and assign overtime; rotate high-demand shifts; integrate recovery time into schedules.
- Satisfaction: Conduct quarterly pulse surveys; visibly respond to feedback.
- Compensation fairness: : Enhance pay structure transparency; provide progress clarity.

- **Tenure management:** Refine onboarding for early-tenure staff; offer mid-career movement to mitigate plateau risk.
- **Commute logistics:** Expand hybrid working and flexitime to suppress DistanceFromHome effects.

3. Institutionalise fairness.

Audit subgroup metrics for Gender and Age; extend to ethnicity, disability, and SES when available. Use parity envelopes to assess disparities. Report that predictions support proactive care, not surveillance (Raji et al., 2020; Mitchell et al., 2019).

4. Combine analytics with judgment.

Combine RF predictions into retention councils where managers triangulate elevated cases with qualitative insights, assuring context-sensitive interventions (Davenport, Harris and Shapiro, 2010).

6.5 Limitations of the Study

This dissertation is rigorous but limited by several boundaries limiting interpretation:

- **Dataset scope and representativeness.** Results are drawn from a single synthetic dataset. While well-designed, it is unable to capture the diversity of drivers across industries and cultures. Attrition factors for retail or healthcare may differ from those in professional services (Shmueli and Koppius, 2011).
- **Cross-sectional and static structure.** The dataset is a snapshot, whereas attrition is a dynamic process. Without longitudinal tracking, the model cannot capture when attrition hazard increase or decrease with time. Hazard timing and unfolding shocks remain unobserved (Angrave et al., 2016).
- **Construct completeness.** Key theoretical drivers—such as sudden external offers, leadership changes, organisational culture, and labour market conditions—were absent. HRIS data often

privilege what is easily measurable over what is most explanatory (Davenport, Harris and Shapiro, 2010).

- **Fairness and bias scope.** Audits were limited to Gender and Age. Factors like ethnicity, disability, socioeconomic status were not available making blind spots in fairness assessment. (Mitchell et al., 2019; Raji et al., 2020).

- **Interpretability and causality.** SHAP explanations increased transparency but not causality. Large feature importance does not confirm that interventions targeting that factor will reduce attrition. Predictions highlight correlations, not causal levers (Shmueli and Koppius, 2011).

- **Governance artefacts** untested in live settings. While Threshold Memos and Model Cards were proposed, they were not stress-tested in organisational practice. Adoption depends on culture, data literacy, and trust, which vary widely across firms (Raji et al., 2020; Mitchell et al., 2019).

6.6 Directions for Future Research

Building on these limitations, future research can expand in five domains:

1. **Cost-sensitive optimisation.** Future studies should apply cost-sensitive RF variants that weight false negatives more heavily, embedding asymmetry directly into model training. Decision-focused evaluation that links predictions to business value can strengthen managerial uptake (Angrave et al., 2016; Davenport, Harris and Shapiro, 2010).

2. **Longitudinal modelling.** Use survival analysis or recurrent-event models with time-varying features (e.g., monthly satisfaction scores, promotion events) to predict attrition hazard windows. This would enable earlier, stage-specific interventions (Shmueli and Koppius, 2011).

3. **Fairness-aware learning.** Expand subgroup audits beyond Gender and Age. Test

governance artefacts and accountability processes (e.g., model cards and internal audits) as integral components of deployment (Mitchell et al., 2019; Raji et al., 2020).

4. **Cross-industry replication.** Compare predictive pipelines by industries (e.g., healthcare, retail, IT) and cultures to test and build sector-specific turnover theory (Angrave et al., 2016).

5. **Human-in-loop uptake studies.** Run A/B field experiments where HR personnel using RF+explanations versus traditional dashboards. Evaluate impacts on retention, intervention success, and staff attitudes. This would advance the given field from model-building to system-embedding (Raji et al., 2020; Mitchell et al., 2019).

6.7 Final Reflections

The most important conclusion of this dissertation is conceptual: attrition is not simply a prediction task but a governed trade-off problem. The challenge is not to find a perfect model but to design a predictive system that is accurate enough, interpretable enough, and governable enough to be trusted by decision makers. Random Forest exemplifies this balance.

For theory, the study reaffirms classical turnover frameworks while reprioritising workload strain. For practice, it offers HR a retention playbook—from rota design to fairness audits—that is grounded in data yet sensitive to organisational realities. For scholarship, it contributes a governance-ready methodology that operationalises threshold policy, calibration, and interpretability. And for me as a researcher, it shifted my understanding from chasing accuracy to stewarding alignment between analytics, ethics, and strategy.

In sum, predictive HR analytics can move beyond being a fashionable buzzword. When embedded responsibly, it becomes a genuine enabler of strategic workforce stability, helping organisations not just to count who leaves but to catch more leavers before they go—fairly, efficiently, and accountably (Davenport, Harris and Shapiro, 2010; Angrave et al., 2016; Shmueli and Koppius, 2011; Mitchell et al., 2019; Raji et al., 2020).

7. References

1. [Hom, P.W., Lee, T.W., Shaw, J.D. and Hausknecht, J.P. \(2017\) 'One hundred years of employee turnover theory and research', *Journal of Applied Psychology*, 102\(3\), pp. 530–545.](#)
2. [Impact of Strategic Human Resource Management Practices on Employee Retention](#)
3. [State of the Global Workplace 2025 \(2025\)](#)
4. [Wright, P.M. and McMahan, G.C. \(1992\) 'Theoretical perspectives for strategic human resource management', *Journal of Management*, 18\(2\), pp. 295–320.](#)
5. [Hausknecht, J.P. and Trevor, C.O. \(2011\) 'Collective turnover at the group, unit, and organizational levels: Evidence, issues, and implications', *Journal of Management*, 37\(1\), pp. 352–388.](#)
6. [Bassi, L. \(2011\) 'Raging debates in HR analytics', *People & Strategy*, 34\(2\), pp. 14–18.](#)
7. [Deloitte \(2017\) *Global Human Capital Trends: Rewriting the rules for the digital age.* Deloitte Insights.](#)
8. [March, J.G. and Simon, H.A. \(1958\) *Organizations*. New York: Wiley.](#)
9. [Price, J.L. \(1977\) *The study of turnover*. Ames: Iowa State University Press.](#)
10. [Price, J.L. \(2001\) 'Reflections on the determinants of voluntary turnover', *International Journal of Manpower*, 22\(7\), pp. 600–624.](#)
11. [Mobley, W.H. \(1977\) 'Intermediate linkages in the relationship between job satisfaction and employee turnover', *Journal of Applied Psychology*, 62\(2\), pp. 237–240.](#)
12. [Porter, L.W. and Steers, R.M. \(1973\) 'Organizational, work, and personal factors in employee turnover and absenteeism', *Psychological Bulletin*, 80\(2\), pp. 151–176.](#)
13. [Herzberg, F., Mausner, B. and Snyderman, B. \(1959\) *The motivation to work*. 2nd edn. New York: Wiley.](#)
14. [Vroom, V.H. \(1964\) *Work and motivation*. New York: Wiley.](#)
15. [Adams, J.S. \(1963\) 'Toward an understanding of inequity', *Journal of Abnormal and Social Psychology*, 67\(5\), pp. 422–436.](#)
16. [Deci, E.L. and Ryan, R.M. \(1985\) *Intrinsic motivation and self-determination in human behavior*. New York: Springer.](#)

17. [Kahn, W.A. \(1990\) 'Psychological conditions of personal engagement and disengagement at work', *Academy of Management Journal*, 33\(4\), pp. 692–724.](#)
18. [Maslach, C. and Leiter, M.P. \(1997\) *The truth about burnout*. San Francisco: Jossey-Bass.](#)
19. [Griffeth, R.W., Hom, P.W. and Gaertner, S. \(2000\) 'A meta-analysis of antecedents and correlates of employee turnover: Update, moderator tests, and research implications for the next millennium', *Journal of Management*, 26\(3\), pp. 463–488.](#)
20. [Mitchell, T.R., Holtom, B.C., Lee, T.W., Sablinski, C.J. and Erez, M. \(2001\) 'Why people stay: Using job embeddedness to predict voluntary turnover', *Academy of Management Journal*, 44\(6\), pp. 1102–1121.](#)
21. [Lee, T.W. and Mitchell, T.R. \(1994\) 'An alternative approach: The unfolding model of voluntary employee turnover', *Academy of Management Review*, 19\(1\), pp. 51–89.](#)
22. [Halbesleben, J.R.B. and Wheeler, A.R. \(2008\) 'The relative roles of engagement and embeddedness in predicting job performance and intention to leave', *Work & Stress*, 22\(3\), pp. 242–256.](#)
23. [Rubenstein, A.L., Eberly, M.B., Lee, T.W. and Mitchell, T.R. \(2018\) 'Surveying the forest: A meta-analysis of voluntary turnover antecedents', *Personnel Psychology*, 71\(1\), pp. 23–65.](#)
24. [Maertz, C.P., Stevens, M.J. and Campion, M.A. \(2007\) 'A turnover model for the 21st century: Integrating social exchange and job embeddedness', *Academy of Management Perspectives*, 21\(1\), pp. 132–148.](#)
25. [Allen, D.G., Bryant, P.C. and Vardaman, J.M. \(2010\) 'Retaining talent: Replacing misconceptions with evidence-based strategies', *Academy of Management Perspectives*, 24\(2\), pp. 48–64.](#)
26. [Batt, R. and Colvin, A.J.S. \(2011\) 'An employment systems approach to turnover: Human resources practices, quits, dismissals, and performance', *Academy of Management Journal*, 54\(4\), pp. 695–717.](#)
27. [Hausknecht, J.P. and Trevor, C.O. \(2011\) 'Collective turnover at the group, unit, and organizational levels: Evidence, issues, and implications', *Journal of Management*, 37\(1\), pp. 352–388.](#)

28. [Angrave, D., Charlwood, A., Kirkpatrick, I., Lawrence, M. and Stuart, M. \(2016\) 'HR and analytics: Why HR is set to fail the big data challenge', *Human Resource Management Journal*, 26\(1\), pp. 1–11.](#)
29. [Davenport, T.H., Harris, J. and Shapiro, J. \(2010\) *Competing on talent analytics*. Boston: Harvard Business Review Press.](#)
30. [Huang, G. and Gartner, J. \(2020\) 'Predicting employee attrition: A machine learning approach', *International Journal of Human Resource Studies*, 10\(2\), pp. 1–20.](#)
31. [Sharma, V. and Mishra, N. \(2021\) 'Predicting employee turnover using machine learning techniques', *International Journal of Advanced Computer Science and Applications*, 12\(5\), pp. 317–325.](#)
32. [Sharma, R. and Ghosh, A. \(2022\) 'A cost-sensitive approach to employee attrition prediction', *International Journal of Manpower*, 43\(4\), pp. 1021–1043.](#)
33. [Floridi, L., Cows, J., Beltrametti, M., Chatila, R., Chazerand, P., Dignum, V., Luetge, C., Madelin, R., Pagallo, U., Rossi, F., Schafer, B., Valcke, P. and Vayena, E. \(2018\) 'AI4People—An ethical framework for a good AI society: Opportunities, risks, principles, and recommendations', *Minds and Machines*, 28\(4\), pp. 689–707.](#)
34. [Barocas, S. and Selbst, A.D. \(2016\) 'Big data's disparate impact', *California Law Review*, 104\(3\), pp. 671–732.](#)
35. [Holstein, K., Wortman Vaughan, J., Daumé, H., Dudik, M. and Wallach, H. \(2019\) 'Improving fairness in machine learning systems: What do industry practitioners need?', *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems*, Glasgow: ACM, pp. 1–16.](#)
36. [Raji, I.D., Smart, A., White, R.N., Mitchell, M., Gebru, T., Hutchinson, B., Smith-Loud, J., Theron, D. and Barnes, P. \(2020\) 'Closing the AI accountability gap: Defining an end-to-end framework for internal algorithmic auditing', *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency \(FAT\)*, Barcelona: ACM, pp. 33–44.](#)
37. [Mitchell, M., Wu, S., Zaldivar, A., Barnes, P., Vasserman, L., Hutchinson, B., Spitzer, E., Raji, I.D. and Gebru, T. \(2019\) 'Model cards for model reporting', *Proceedings of the 2019 Conference on Fairness, Accountability, and Transparency \(FAT\)*, Atlanta: ACM, pp. 220–229.](#)

38. [Kleinberg, J., Mullainathan, S. and Raghavan, M. \(2017\) 'Inherent trade-offs in the fair determination of risk scores', *Proceedings of the 8th Innovations in Theoretical Computer Science Conference \(ITCS\)*, Berkeley: ACM, pp. 43–51.](#)
39. [Mehrabi, N., Morstatter, F., Saxena, N., Lerman, K. and Galstyan, A. \(2021\) 'A survey on bias and fairness in machine learning', *ACM Computing Surveys*, 54\(6\), pp. 1–35.](#)
40. [Leicht-Deobald, U., Busch, T., Schank, C., Weibel, A., Schafheitle, S., Wildhaber, I. and Kasper, G. \(2019\) 'The challenges of algorithm-based HR decision-making for personal integrity', *Journal of Business Ethics*, 160\(2\), pp. 377–392.](#)
41. [Lepri, B., Oliver, N., Letouzé, E., Pentland, A. and Vinck, P. \(2018\) 'Fair, transparent, and accountable algorithmic decision-making processes', *Philosophy & Technology*, 31\(4\), pp. 611–627.](#)
42. [Saunders, M., Lewis, P. and Thornhill, A. \(2019\) *Research methods for business students*. 8th edn. Harlow: Pearson.](#)
43. [Bryman, A. and Bell, E. \(2015\) *Business research methods*. 4th edn. Oxford: Oxford University Press.](#)
44. [Shmueli, G. and Koppius, O.R. \(2011\) 'Predictive analytics in IS research', *MIS Quarterly*, 35\(3\), pp. 553–572.](#)
45. [Tranfield, D., Denyer, D. and Smart, P. \(2003\) 'Systematic review methodology', *British Journal of Management*, 14\(3\), pp. 207–222.](#)
46. [Kuhn, M. and Johnson, K. \(2013\) *Applied predictive modeling*. New York: Springer.](#)
47. [Kuhn, M. and Silge, J. \(2017\) *Tidy modeling with R*. Sebastopol: O'Reilly.](#)
48. [Efron, B. and Tibshirani, R.J. \(1994\) *An introduction to the bootstrap*. New York: Chapman & Hall.](#)
49. [Kohavi, R. \(1995\) 'Cross-validation accuracy estimation', in *Proceedings of the International Joint Conference on Artificial Intelligence \(IJCAI\)*, pp. 1137–1145.](#)
50. [Rahm, E. and Do, H.H. \(2000\) 'Data cleaning: Problems and approaches', *IEEE Data Engineering Bulletin*, 23\(4\), pp. 3–13.](#)
51. [IBM \(2019\) *IBM HR Analytics Dataset*. Available at: <https://www.ibm.com/> \(Accessed: 8 September 2025\).](#)

52. [Breiman, L. \(2001\) 'Random forests', *Machine Learning*, 45\(1\), pp. 5–32.](#)
53. [Breiman, L., Friedman, J.H., Olshen, R.A. and Stone, C.J. \(1984\) *Classification and regression trees*. Belmont, CA: Wadsworth.](#)
54. [Chen, T. and Guestrin, C. \(2016\) 'XGBoost: A scalable tree boosting system', in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining \(KDD '16\)*. New York: ACM, pp. 785–794.](#)
55. [Cortes, C. and Vapnik, V. \(1995\) 'Support-vector networks', *Machine Learning*, 20\(3\), pp. 273–297.](#)
56. [Chawla, N.V., Bowyer, K.W., Hall, L.O. and Kegelmeyer, W.P. \(2002\) 'SMOTE: Synthetic Minority Over-sampling Technique', *Journal of Artificial Intelligence Research*, 16, pp. 321–357.](#)
57. [Lemaître, G., Nogueira, F. and Aridas, C.K. \(2017\) 'Imbalanced-learn: A Python toolbox to tackle the curse of imbalanced datasets in machine learning', *Journal of Machine Learning Research*, 18\(17\), pp. 1–5.](#)
58. [Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O. et al. \(2011\) 'Scikit-learn: Machine learning in Python', *Journal of Machine Learning Research*, 12, pp. 2825–2830.](#)
59. [Niculescu-Mizil, A. and Caruana, R. \(2005\) 'Predicting good probabilities with supervised learning', in *Proceedings of the 22nd International Conference on Machine Learning \(ICML '05\)*. New York: ACM, pp. 625–632.](#)
60. [Elkan, C. \(2001\) 'The foundations of cost-sensitive learning', in *Proceedings of the 17th International Joint Conference on Artificial Intelligence \(IJCAI '01\)*. San Francisco, CA: Morgan Kaufmann, pp. 973–978.](#)
61. [Saito, T. and Rehmsmeier, M. \(2015\) 'The precision–recall plot is more informative than the ROC plot when evaluating binary classifiers on imbalanced datasets', *PLOS ONE*, 10\(3\), e0118432.](#)

